

What is Research in Programming Languages?

Prof. Dr. Guido Salvaneschi

School of Computer Science



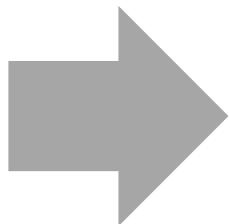
University of St. Gallen

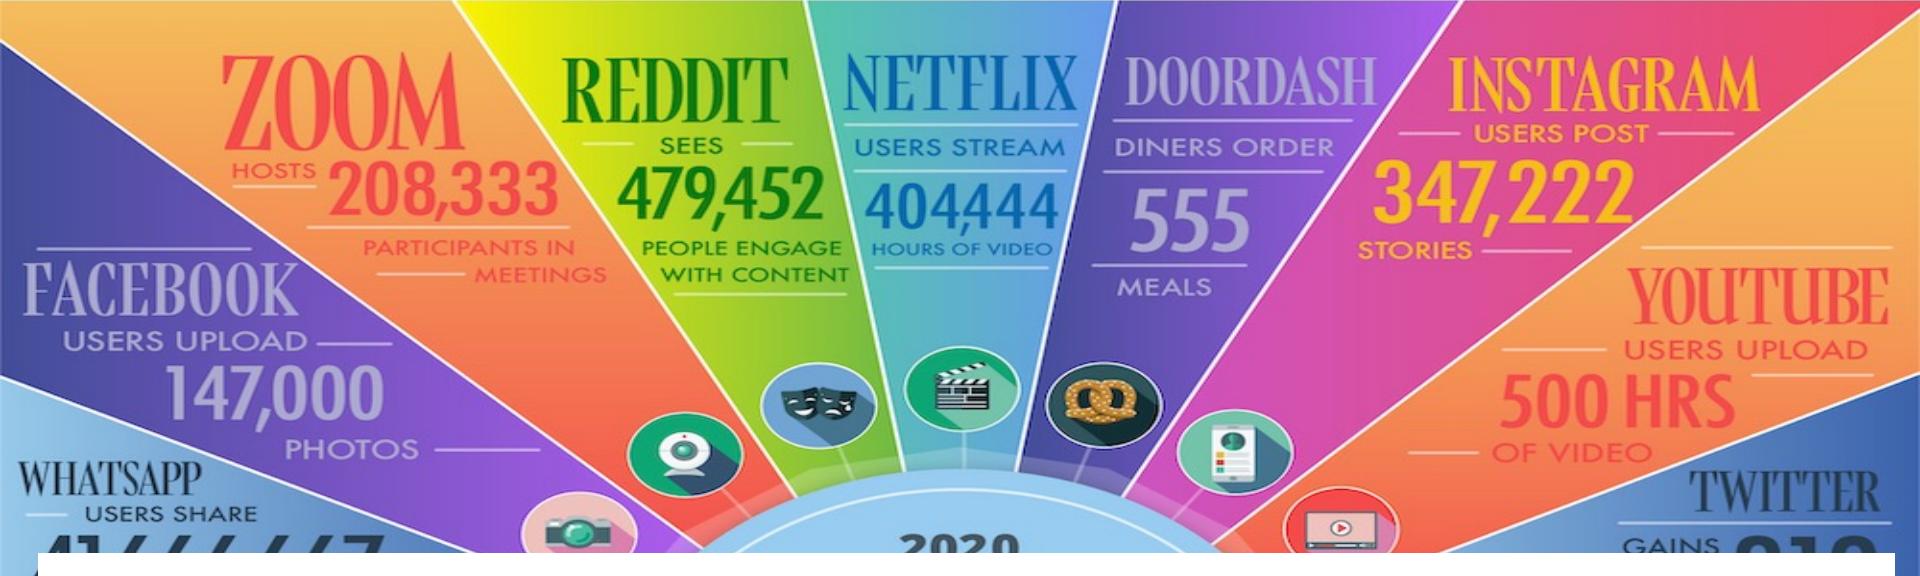
The World as we Knew it

```
use AdventureWorksDW2008R2

SELECT e.DepartmentName, COUNT(*)
FROM [dbo].[DimEmployee] AS e
WHERE e.Gender = 'F' and e.SickLeaveCount > 10
GROUP BY e.DepartmentName
HAVING COUNT(*) > 3
ORDER BY EmployeeCount DESC
```

	DepartmentName	EmployeeCount
1	Production	29
2	Purchasing	6
3	Finance	5
4	Information Services	4





2 years = 90% of data ever generated

By 2020: 1.7MB per second by each person



Real Time Processing



Fraud Detection



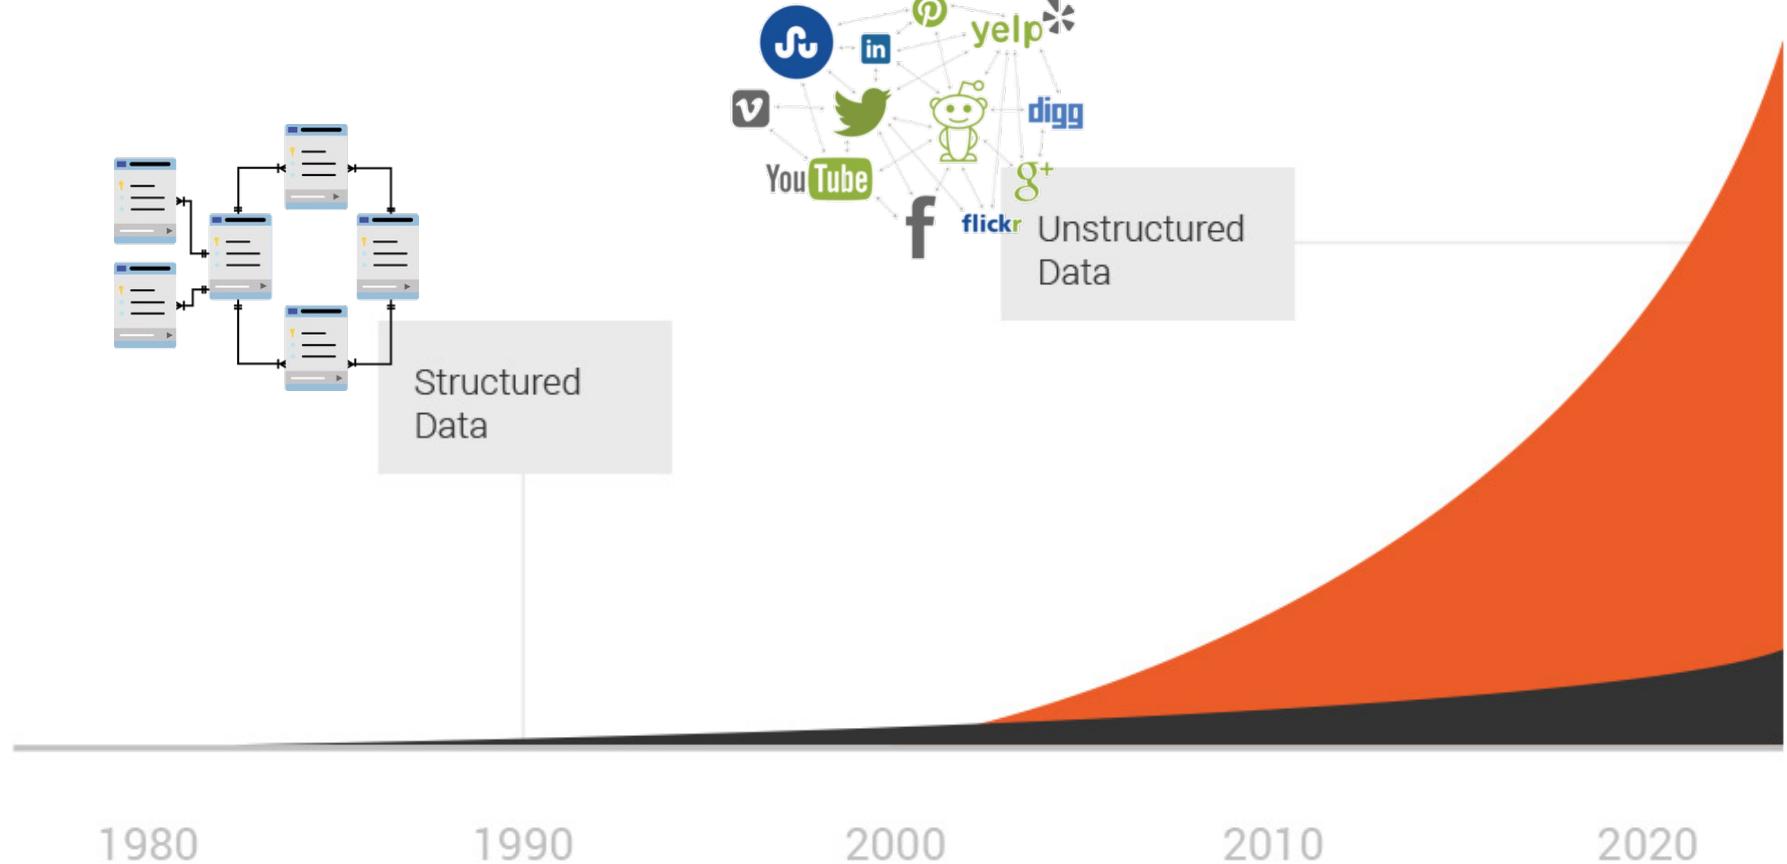
Real Time Business
Intelligence



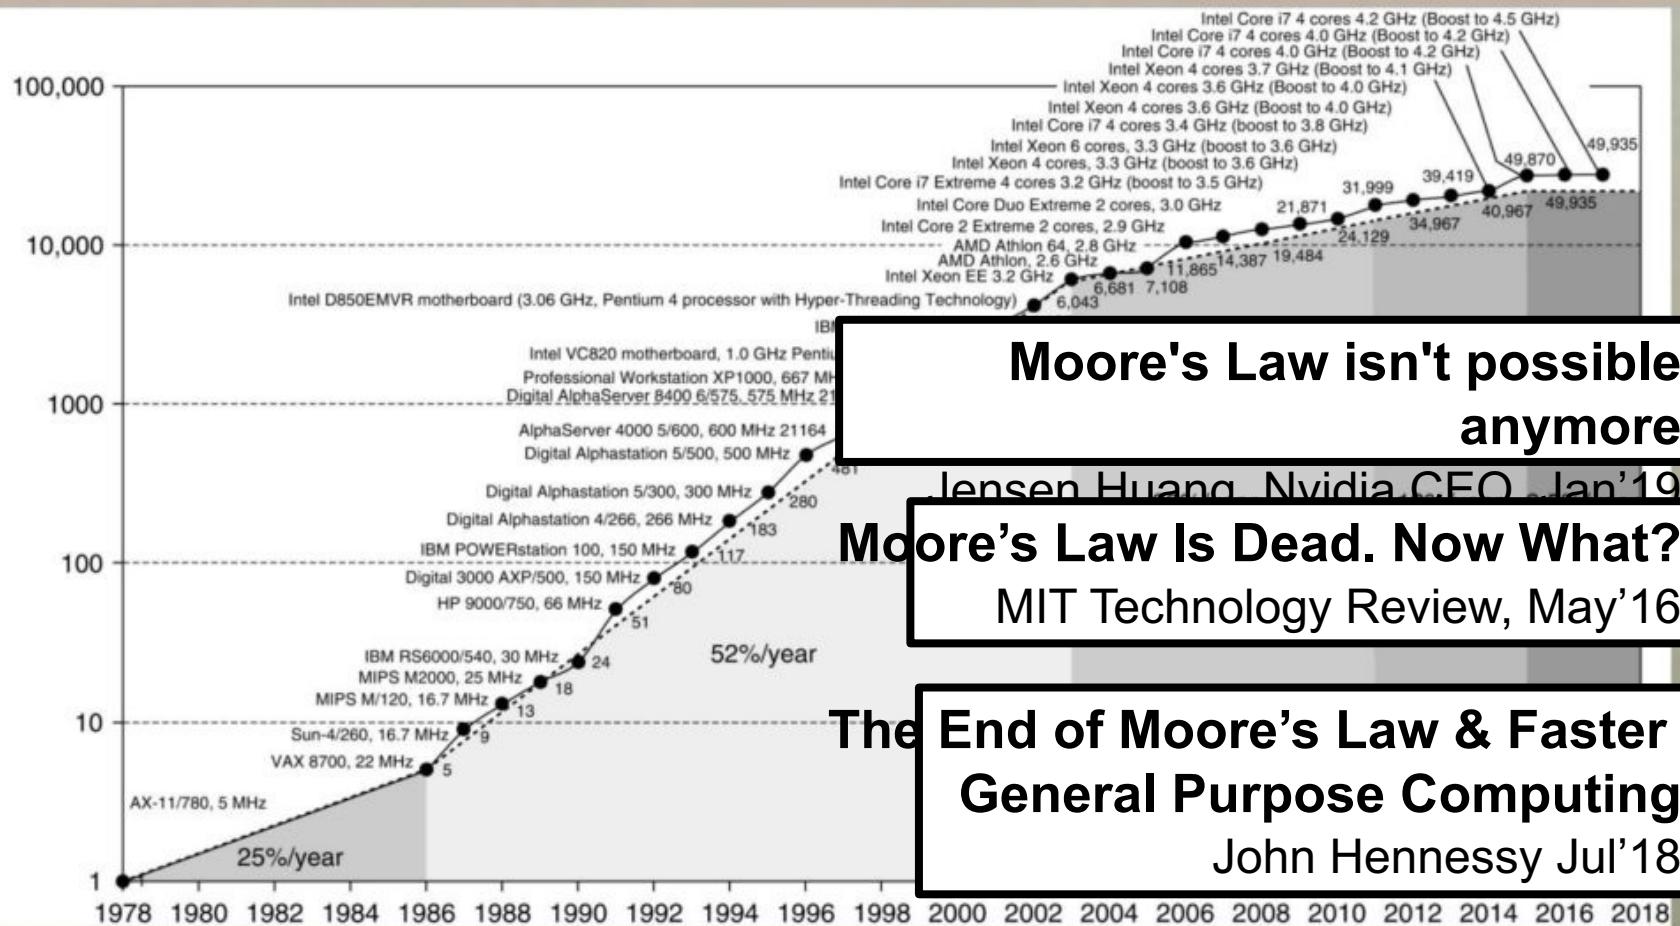
Cloud Monitoring

Structured and Unstructured Data

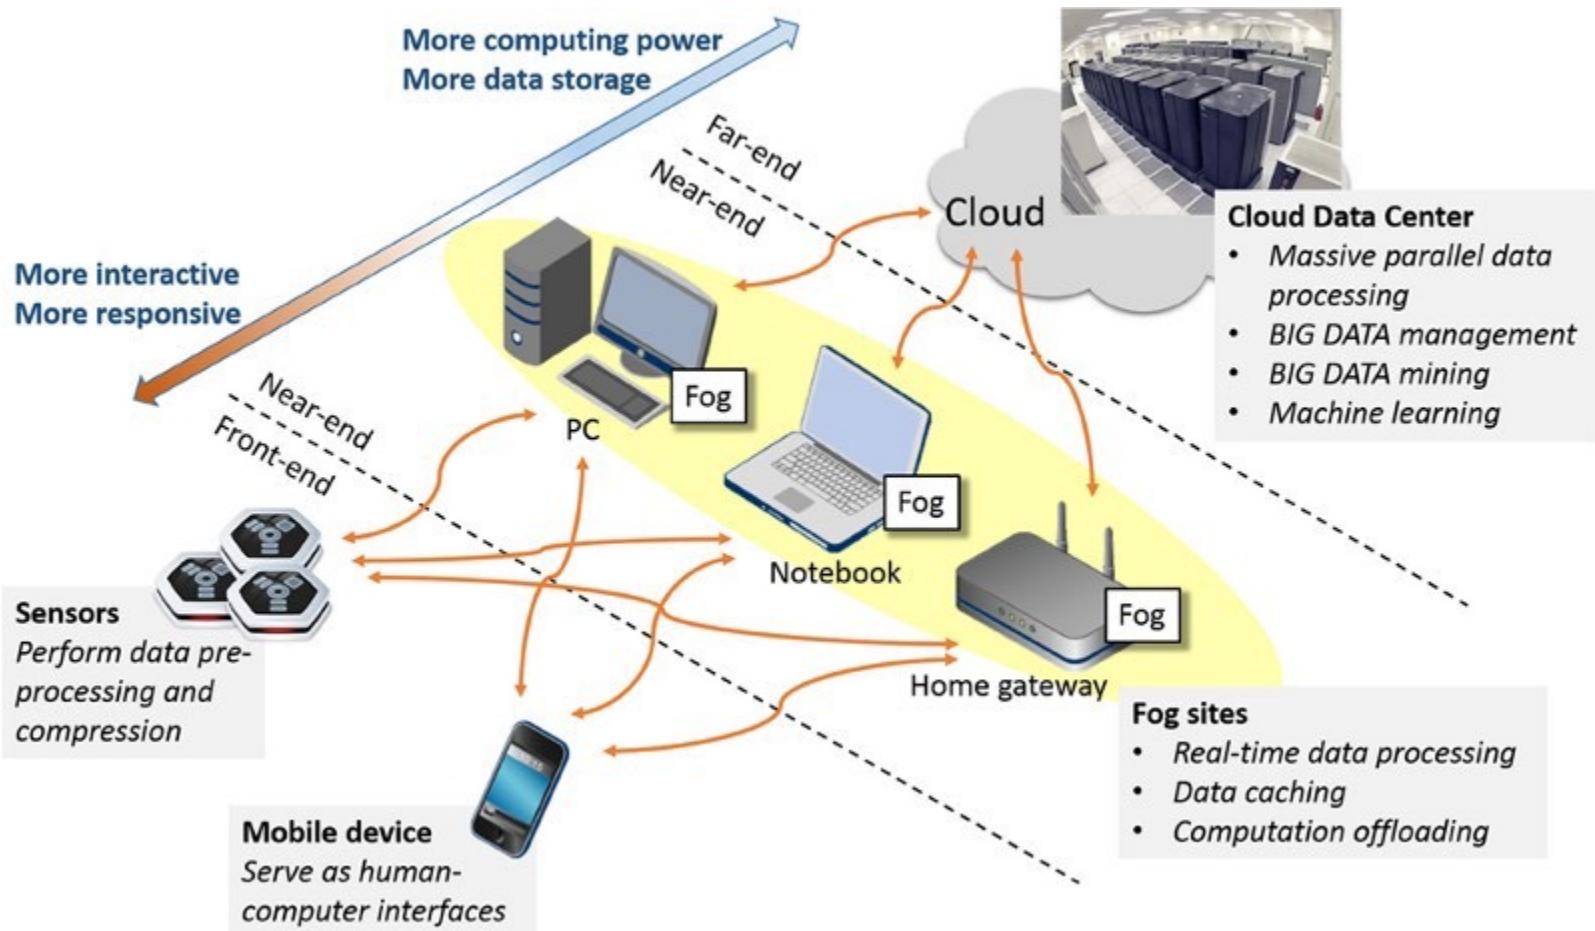
IDC and EMC project that data will grow to 40 ZB by 2020



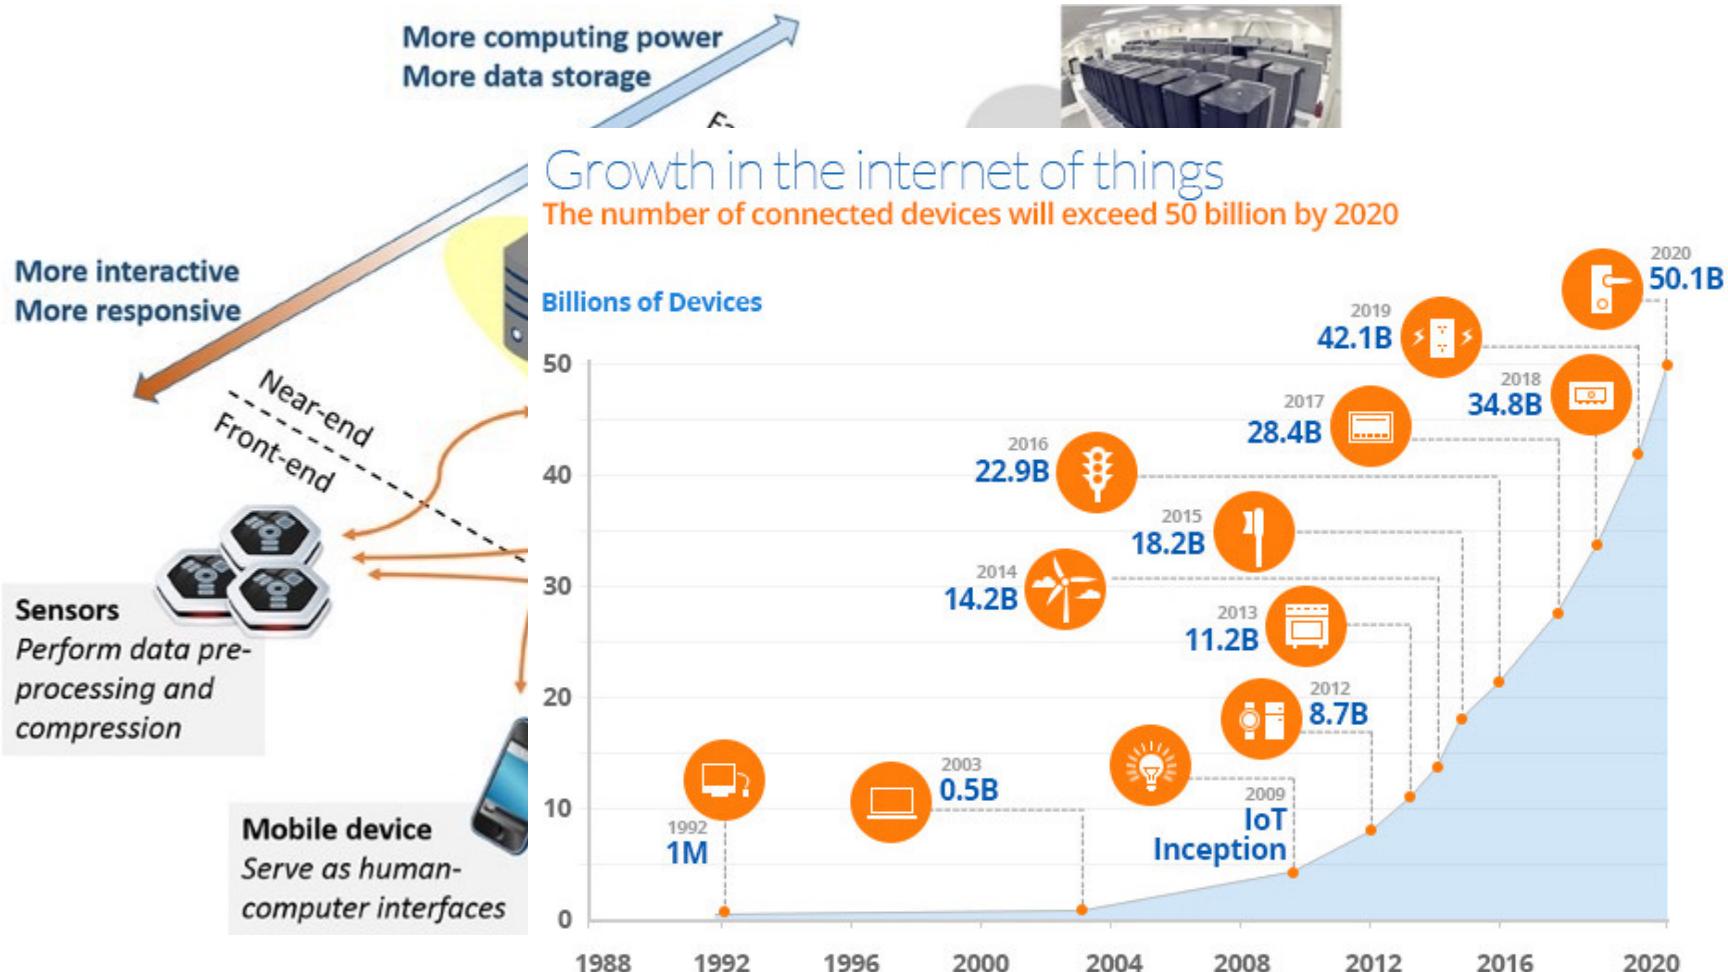
UNIPROCESSOR PERFORMANCE (SINGLE CORE)

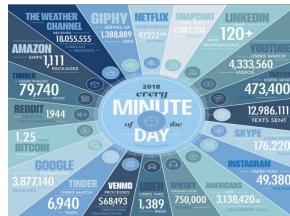


Processing at the Edge

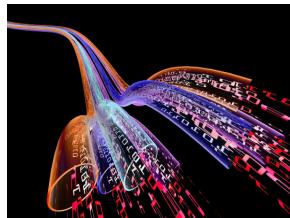


Processing at the Edge

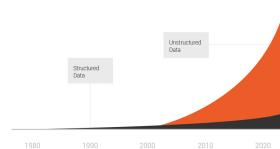




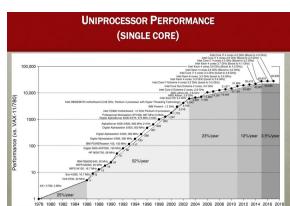
Big Data



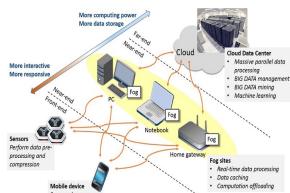
Structured and Unstructured Data
IDC and EMC project that data will grow to 40 ZB by 2020



Real Time Requirements



No Moore's Law



Edge

DATA-INTENSIVE DISTRIBUTED APPLICATIONS

SCALABLE

LOW LATENCY



EVENT BASED

DISTRIBUTED

HETEROGENEOUS

SOFTWARE DESIGN



Apache Flink® – Stateful Computations over Data Streams

What is Apache Flink?

Use Cases

Powered By

FAQ

Downloads

Tutorials

Documentation

Getting Help

Flink Blog

Community & Project Info

Roadmap

How to Contribute

Flink on GitHub

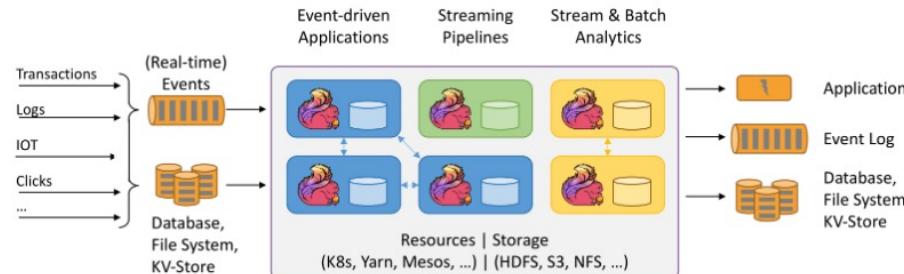
中文版

@ApacheFlink

Plan Visualizer

Apache Software Foundation

License Security
Donate Thanks

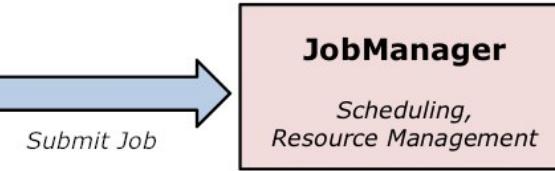
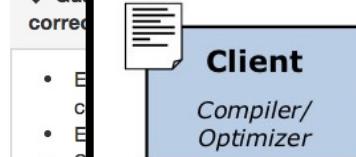


All streaming use cases

- Event-driven Applications
- Stream & Batch Analytics
- Data Pipelines & ETL

[Learn more](#)

Program



Operational Focus

- Flexible deployment
- High-availability setup
- Savepoints

[Learn more](#)

Scalability

- Stateful
- Stateless
- Incremental

[Learn more](#)

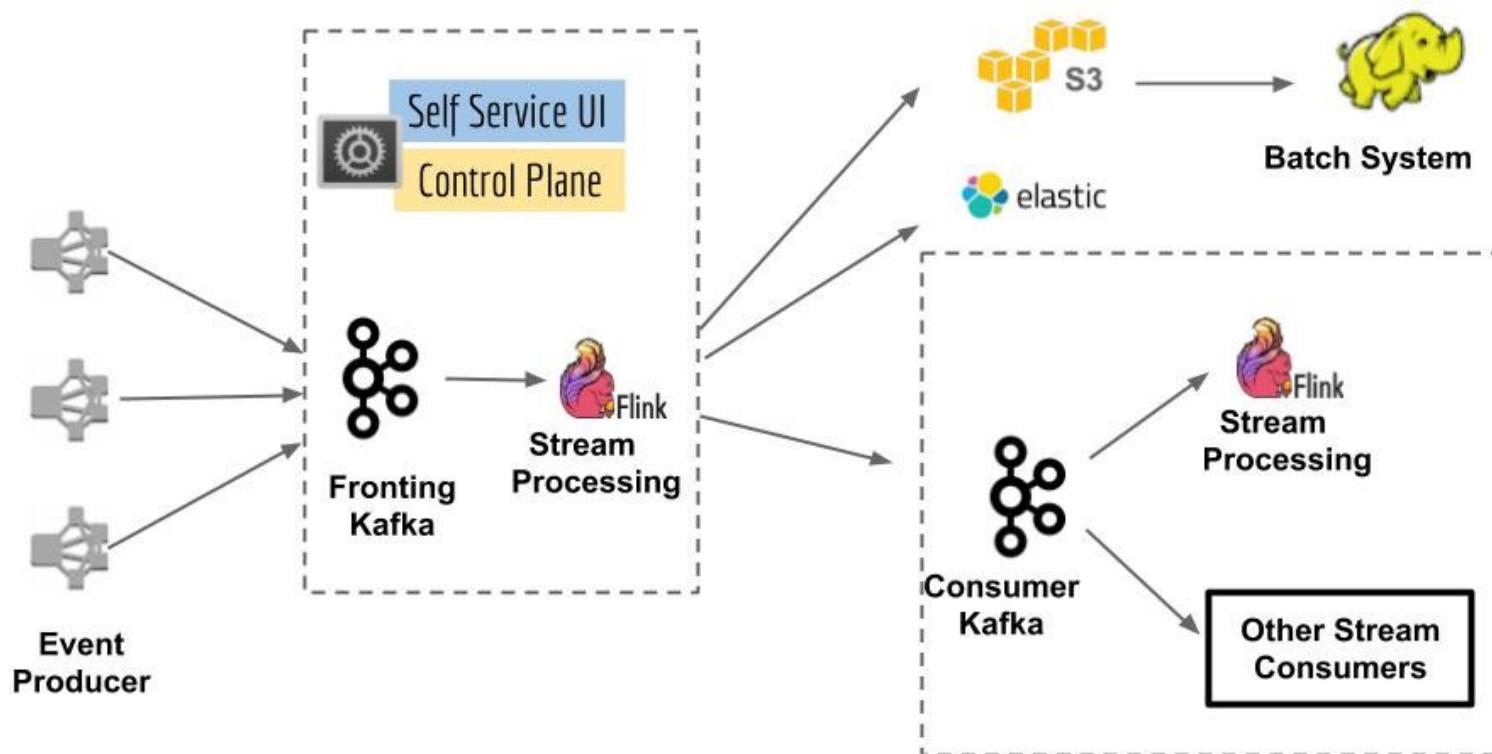


Powered by Flink





Inca: Netflix's real time message tracing system



FRAMEWORKS

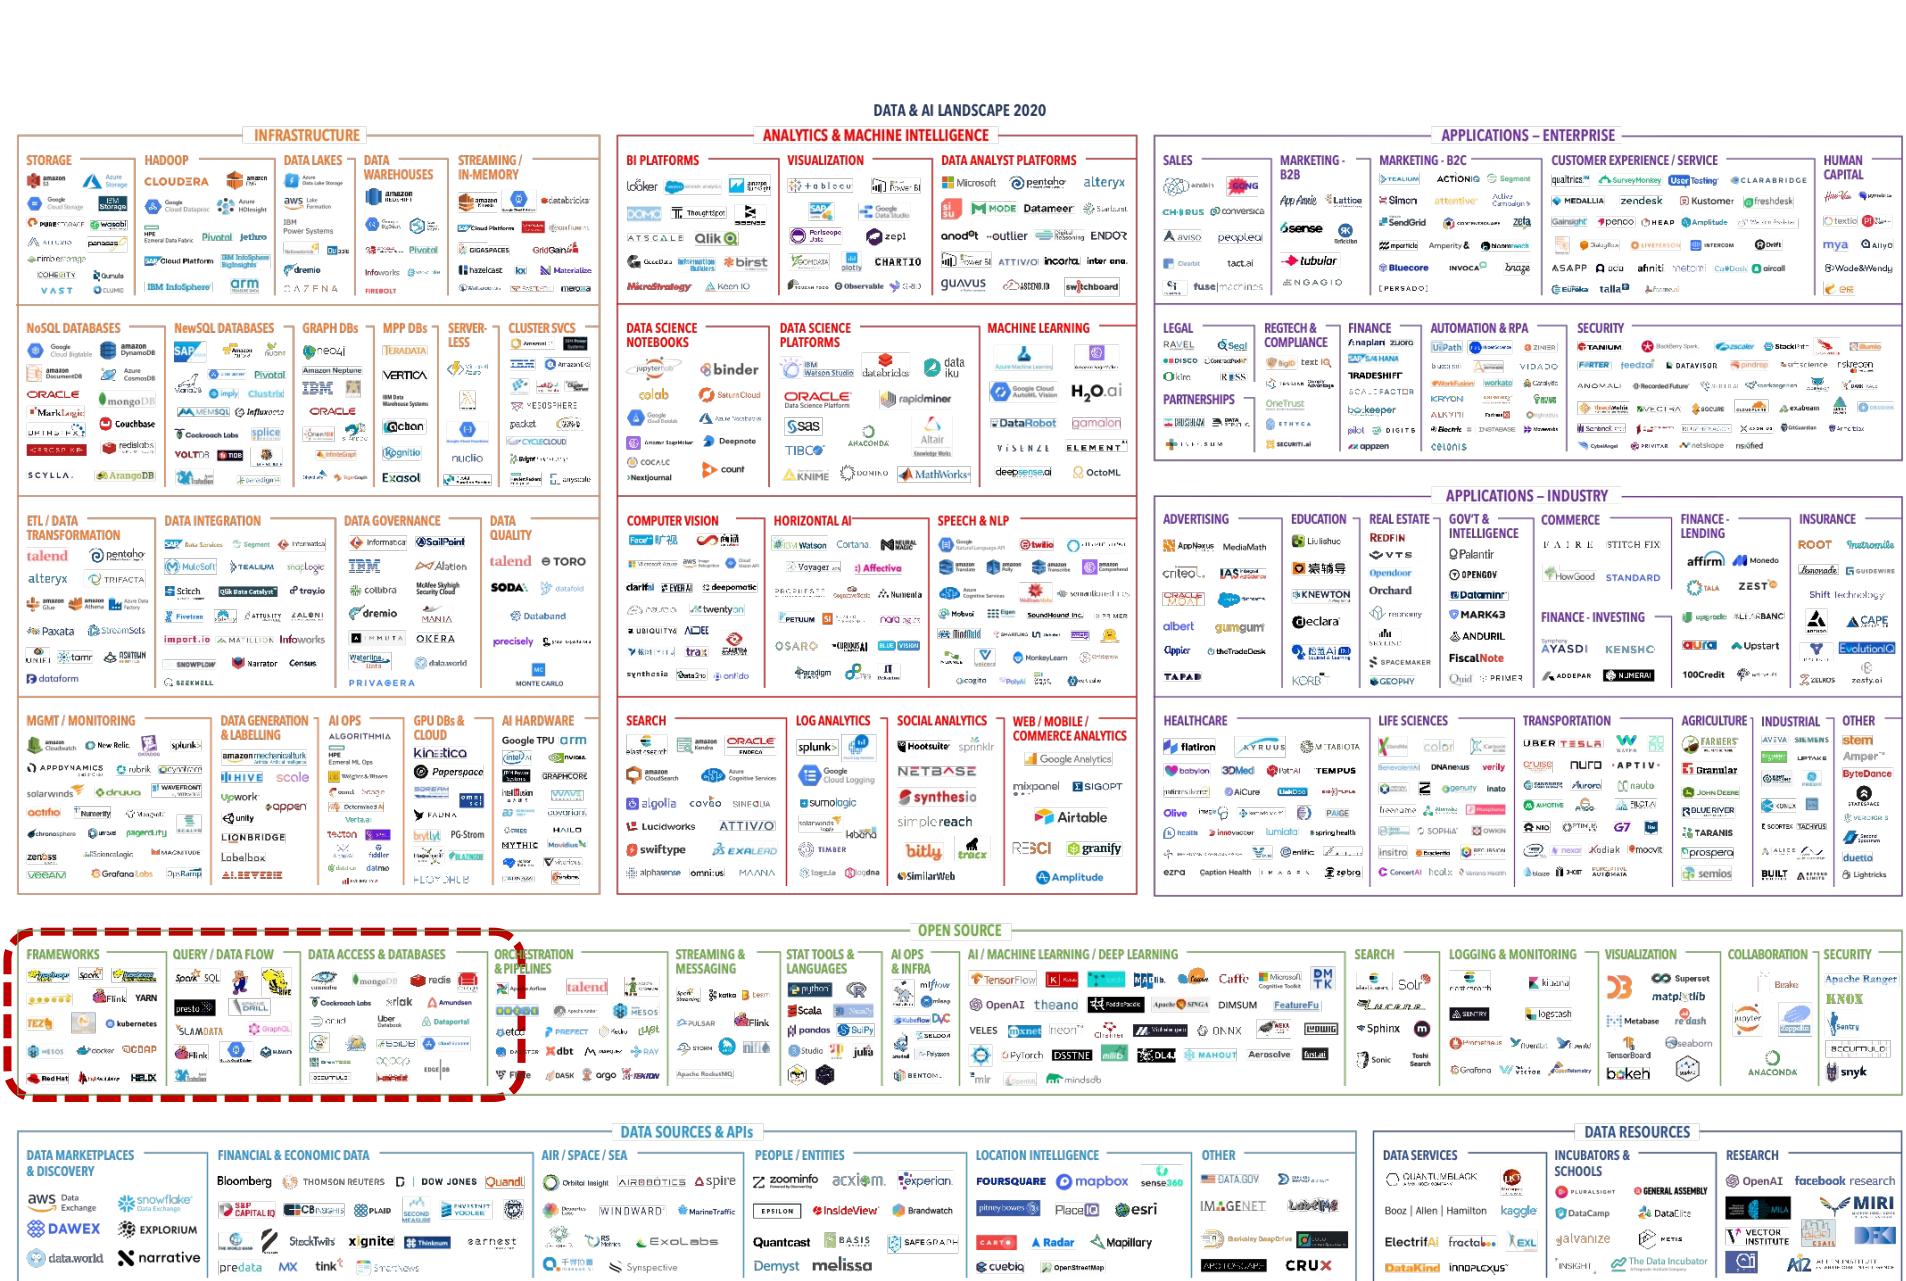


QUERY / DATA FLOW

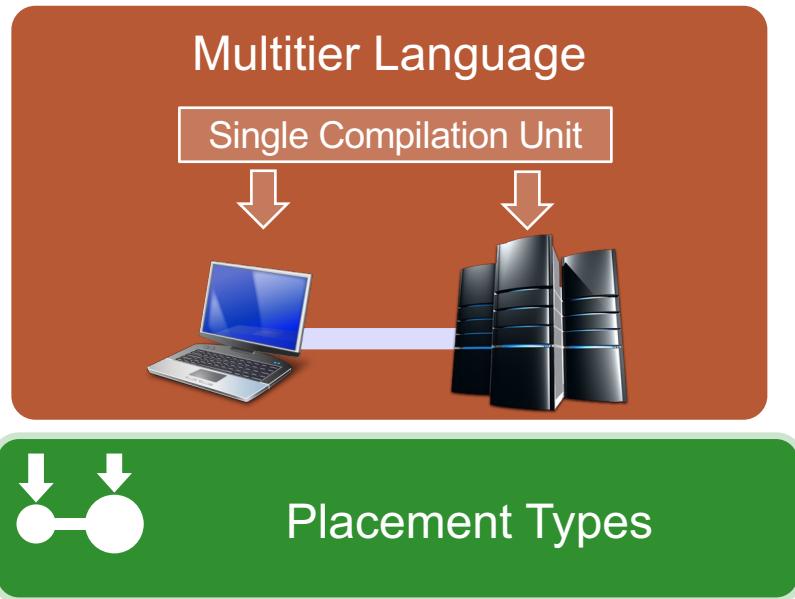


DATA ACCESS & DATABASES





ScalaLoci Programming Framework



www.scala-loci.github.io



[P.Weisenbureger, M.Koehler, G.Salvanesci, **Distributed System Development with ScalaLoci**, OOPSLA'18]

[P.Weisenbureger, G.Salvanesci, **Multitier Modules**, ECOOP'19]

Placement Types

```
trait Registry extends Peer  
trait Node extends Peer
```

Peers

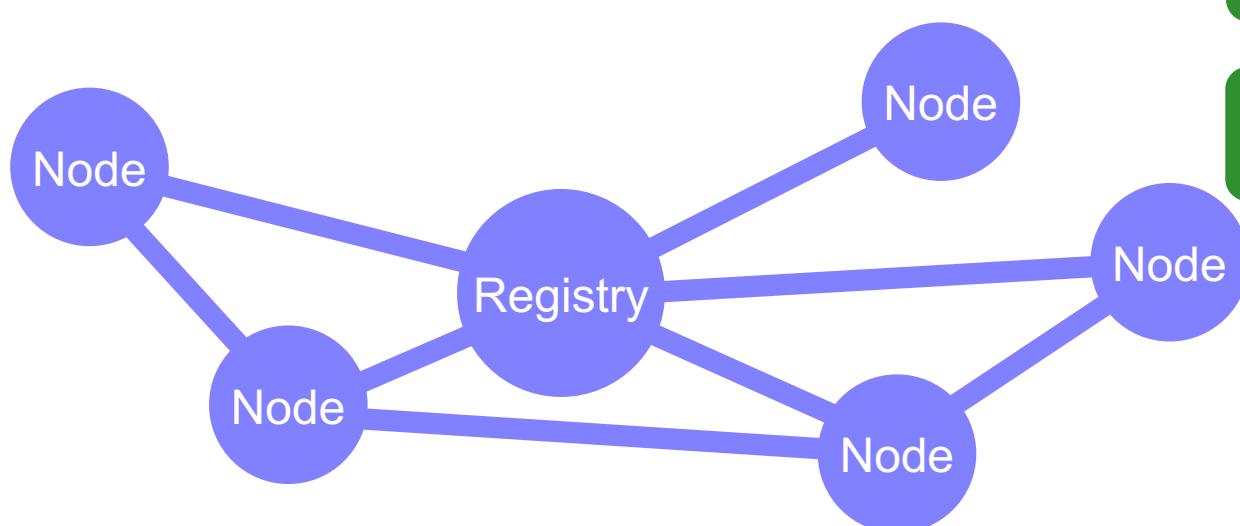
```
val message: Event[String] on Registry
```

Placement Types

Architecture

```
trait Registry extends Peer { type Tie = Multiple[Node] }
```

```
trait Node extends Peer { type Tie = Single[Registry] with Multiple[Node]
```

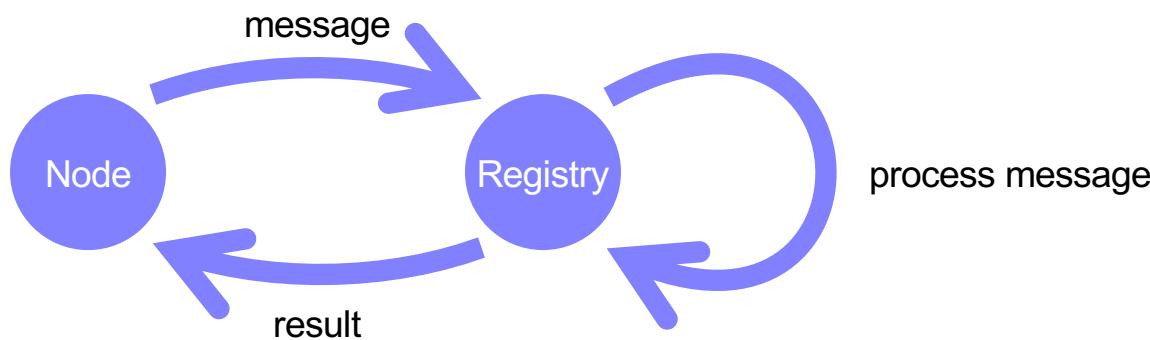


Architecture Specification
through Peer Types

Architecture-Based
Remote Access

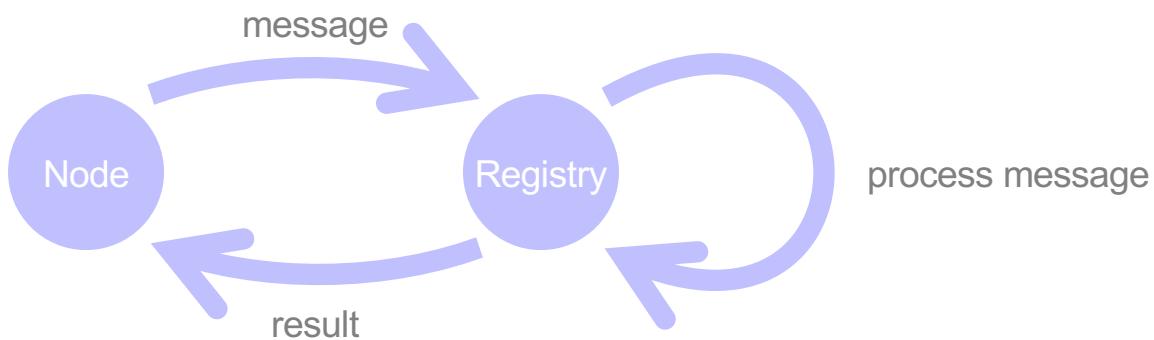
Data Flow

```
val message = Event[String]()
val result = message map processMessage
val ui = new UI(result)
```



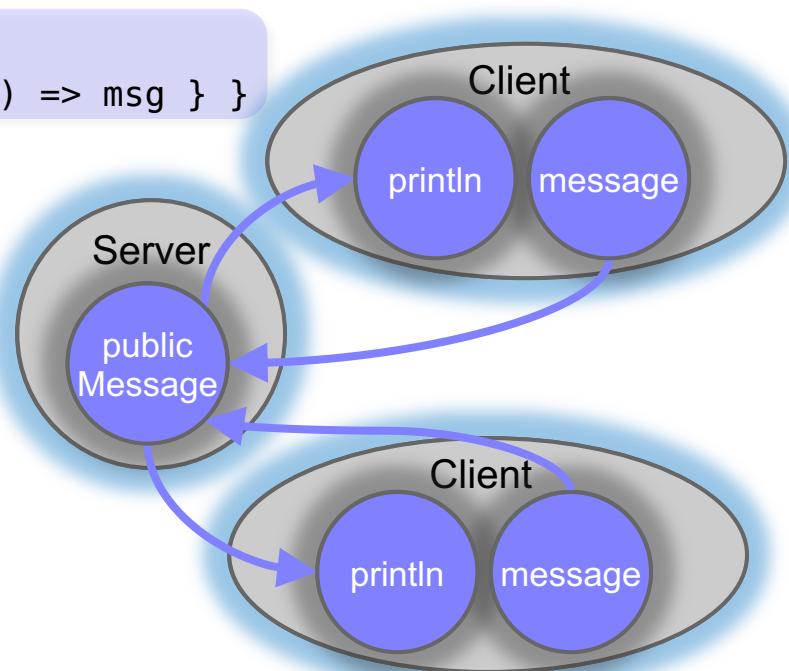
Distributed Data Flow

```
val message: Event[String] on Node = placed[Node] { Event[String]() }
val result = placed[Registry] { message.asLocal map processMessage }
val ui = placed[Node] { new UI(result.asLocal) }
```



Complete Distributed Chat

```
@multitier object Chat {  
  trait Server extends Peer { type Tie = Multiple[Client] }  
  trait Client extends Peer { type Tie = Single[Server] }  
  
  val message = placed[Client] { Evt[String] }  
  
  val publicMessage = placed[Server] {  
    message.asLocalFromAllSeq map { case (_, msg) => msg } }  
  
  placed[Client].main {  
    publicMessage.asLocal observe println  
    for (line <- io.Source.stdin.getLines)  
      message fire line } }
```



Porting to Distribution

Local

```

val ballSize = 20
val maxy = 30
val maxx = 400
val leftPos = 30
val rightPos = 770
val initPosition = Point(400, 200)
val initSpeed = Point(10, 8)

val ball: Signal[Point] = tick.fold(initPosition) {
  (ball, _) => ball + speed.get
}

val areas = {
  val racket = Seq(
    Signal[UI.mousePosition].y,
    Signal[UI.mousePosition].y + 100
  )
  val leftRacket = Racket(leftRacketPos, racket(0))
  val rightRacket = Racket(rightRacketPos, racket(1))
  val collisionRacket = CollisionRacket(leftRacket)
  val collisionRacket = CollisionRacket(rightRacket)
  Signal[Rects] (rects map { _area })
}

val leftWall = ball.changed && (-x < 0)
val rightWall = ball.changed && (_x > maxx)

val yBounce = {
  val ballInRacket = Signal[Areas] (areas exists { _contains ball() })
  val collisionRacket = ballInRacket.changed && true
  val yBounce = ballInRacket.changed && collisionRacket
  val yBounce = ball.changed &&
    (ball.y <= ball.y || ball.y > maxy)
}

val speed = {
  val x = {x: Double => toggle (initSpeed.x, -initSpeed.x)}
  val y = {y: Double => toggle (initSpeed.y, -initSpeed.y)}
  Signal[Point](Point(x(), y()))
}

val score = {
  val leftPoints = rightWall.ite(0, { -1 })
  val rightPoints = leftWall.ite(0, { -1 })
  Signal[LeftPoints](leftPoints() + ":" + rightPoints())
}

val ui = new UI(areas, ball, score)

```

ScalaLoci

```

trait Server extends Actor
{
    val batSize = 20
    val maxX = 800
    val maxY = 400
    val leftPos = 100
    val rightPos = 770
    val initPosition = Pos(leftPos, 200)
    val initSpeed = 10
}

val clientMouse = @labeled[Player]
Signal[UI.mouses]

val isPlaying = @labeled[Player]
Signal[Client.Client]

val ball = Signal[Point]
tick.fold(instantPosition) {
    if (isPlaying.get)
        isPlaying.get
}

val players = @labeled[Player]
Remote[Client.Client]
case left :: right
    case _ :: Seq(left)
    case _ :: Seq(right)

val areas = @labeled[Player]
Client.Client
instantPosition() {
    val left = Point(leftPos, 200)
    val right = Point(rightPos, 200)
    val rackets = List[Signal[Player.Racket]](left, right)
    rackets.map(_.get)
}

val leftWall = @labeled[Player]
val rightWall = @labeled[Player]

val xBounce = @labeled[Player]
val yBounce = @labeled[Player]
val ballBounces = Signal[Point]
leftWall || rightWall
leftWall || rightWall
yBounce = placed(xBounce.get, yBounce.get)

val ball = Signal[Point]
ball.bounce() {
    ball = ball.bounce(yBounce.get)
}

val speed = @labeled[Player]
val x = "bounce" toggle
val y = "bounce" toggle
Signal[Point(x, y)]

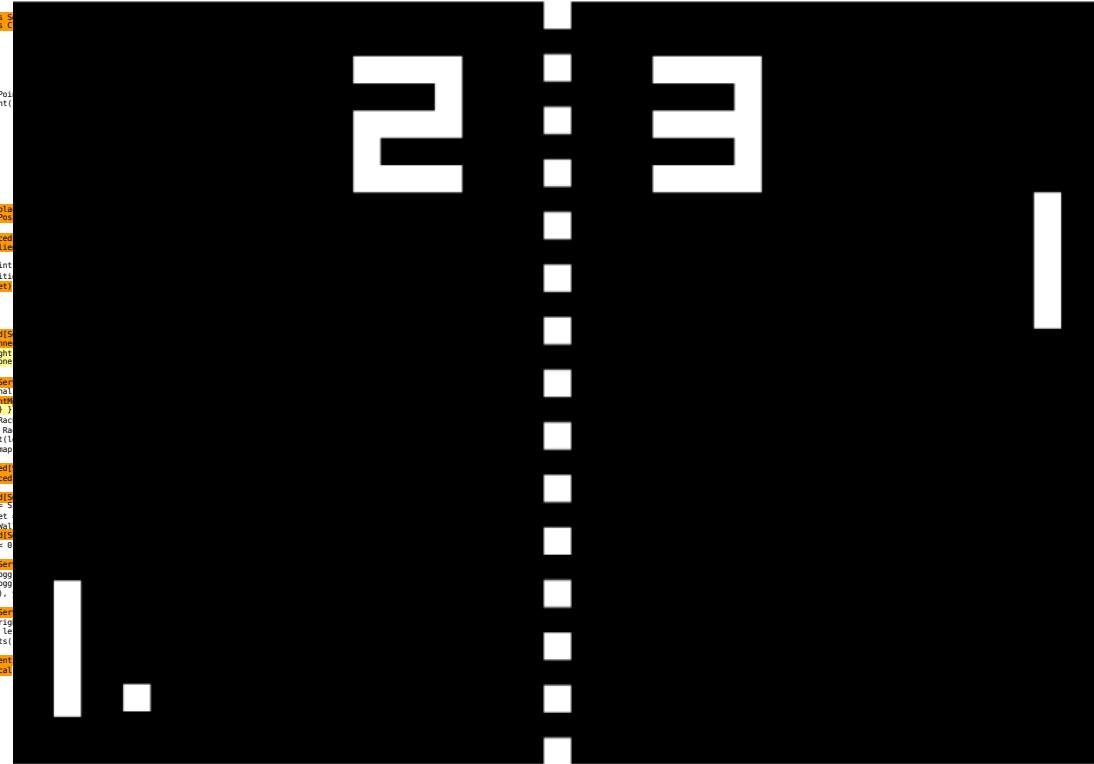
val score = @labeled[Player]
val leftPoints = rightPoints = 0
Signal[Point(x, y)]

val ui = @labeled[Client.Client]

```

Akka

RMI



```
case UpdateScore(score) => this.score set score  
server ! AddPlayer
```

```

clients.foreach { _updateScore(score) }

private trait Client {
  def updateScore(area: List[Area]): Unit
  def updateScore(score: String): Unit
}

class ClientImpl(server: Server) extends Client {
  val self = makeStub[Client](this)

  val areas = Var(List.empty[Area])
  val score = Var("0" * 9)

  UI.mouseMove.observe { pos =>
    server.mouseChanged(self, pos.y)
  }

  val ui = new UI(area, ball, score)

  def updateScore(area: List[Area]) = synchronized { this.areas = areas }
  def updateBall(ball: Point) = synchronized { this.ball = ball }
  def updateScore(score: String) = synchronized { this.score() = score }

  server.addPlayer(self)
}

```



Flink

```
class TaskManagerGateway {
  def handleExceptionFromJobManager(instanceId: InstanceID, cause: Exception,
  ) = {
    m! (instanceId, instanceId, cause)
  }
  def handleCluster(applicationStatus: ApplicationStatus, message: String,
  ) = {
    m! (applicationStatus, applicationStatus, message)
  }
  def handleStackTrace(StackTrace: StackTrace) = {
    m! (StackTrace, StackTrace)
  }
  def handleTask(trid: TaskDeploymentID, mgr: ActorRef) = {
    m! (trid, (trid, mgr))
  }
  def handleTask(executionAttemptID: ExecutionAttemptID, mgr: ActorRef) = {
    m! (executionAttemptID, (executionAttemptID, mgr))
  }
  def handleTask(executionAttemptID: ExecutionAttemptID, mgr: ActorRef) = {
    m! (executionAttemptID, (executionAttemptID, mgr))
  }
  def handleTask(partitionID: PartitionID, mgr: ActorRef) = {
    m! (partitionID, (partitionID, mgr))
  }
  def updatePartitions(executionAttemptID: ExecutionAttemptID, confs: Map[PartitionID, PartitionConf], rebalance: Boolean) = {
    m! (executionAttemptID, (executionAttemptID, confs, rebalance))
  }
  def handlePartition(executionAttemptID: ExecutionAttemptID, jobID: JobID) = {
    m! (executionAttemptID, (jobID, executionAttemptID))
  }
  def handlePartition(executionAttemptID: ExecutionAttemptID, jobID: JobID) = {
    m! (executionAttemptID, (jobID, executionAttemptID))
  }
  def handleCheckpointCompleted(executionAttemptID: ExecutionAttemptID, jobID: JobID, checkpointID: CheckpointID, checkpointType: CheckpointType) = {
    m! (executionAttemptID, (jobID, executionAttemptID, checkpointID, checkpointType))
  }
  def triggerCheckpoint(int(executionAttemptID: ExecutionAttemptID, jobID: JobID, checkpointID: CheckpointID, checkpointType: CheckpointType, times: Int, checkpointOptions: CheckpointOptions)) = {
    m! (executionAttemptID, (jobID, executionAttemptID, checkpointID, checkpointType, times, checkpointOptions))
  }
  def handleTypeRequest[T](jobTypeRequest: JobTypeRequest, mgr: ActorRef) = {
    m! (jobTypeRequest.jobID, (jobTypeRequest, mgr))
  }
  def handleTypeRequest[T](jobTypeRequest: JobTypeRequest, mapTo: BlockKey) = {
    m! (jobTypeRequest.jobID, (jobTypeRequest, mapTo))
  }
}
```



Crosscutting functionality separated among compilation units

```
class JobManager extends Actor {
  ...
  def cluster(applicationStatus: ApplicationStatus, message: String) {
    ...
  }
  ...
  def TaskManager extends Actor {
    ...
    def TaskManager(jobManager: JobManager, actorRef: ActorRef) = {
      ...
    }
    ...
  }
  ...
}
```



Flink



Developers are **not** forced to modularize **along network boundaries**

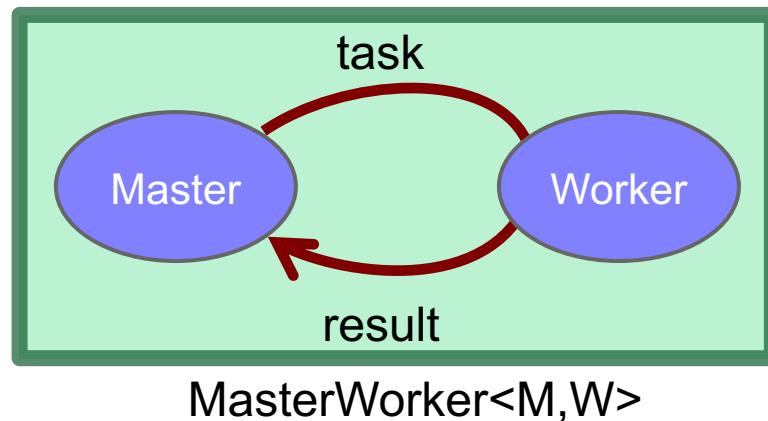
```
class JobManager extends Actor {
  ...
  def cluster(applicationStatus: ApplicationStatus, message: String) {
    ...
  }
  ...
  def TaskManager(jobManager: JobManager, actorRef: ActorRef) = {
    ...
  }
  ...
}
```



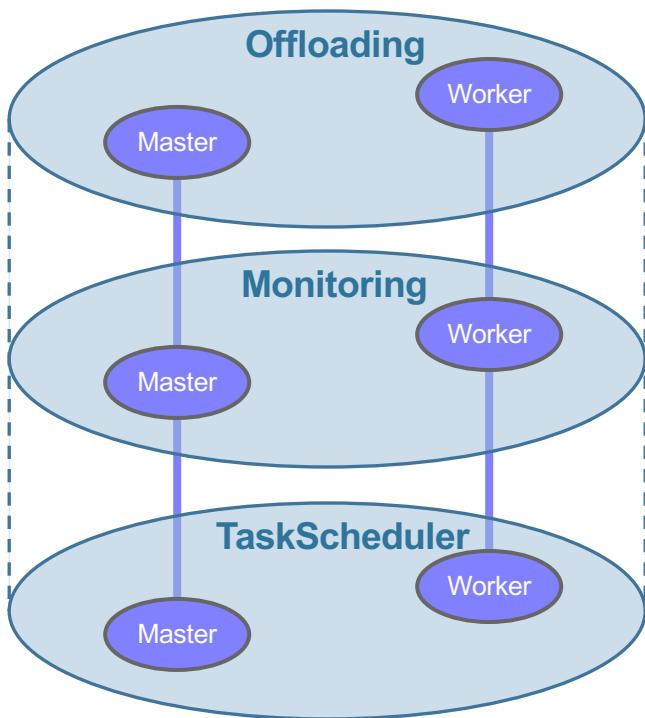
Modelling Distributed Software: Multitier Modules

Distributed functionality = Module

Composing modules = Composing subsystems



Stacking Multitier Modules



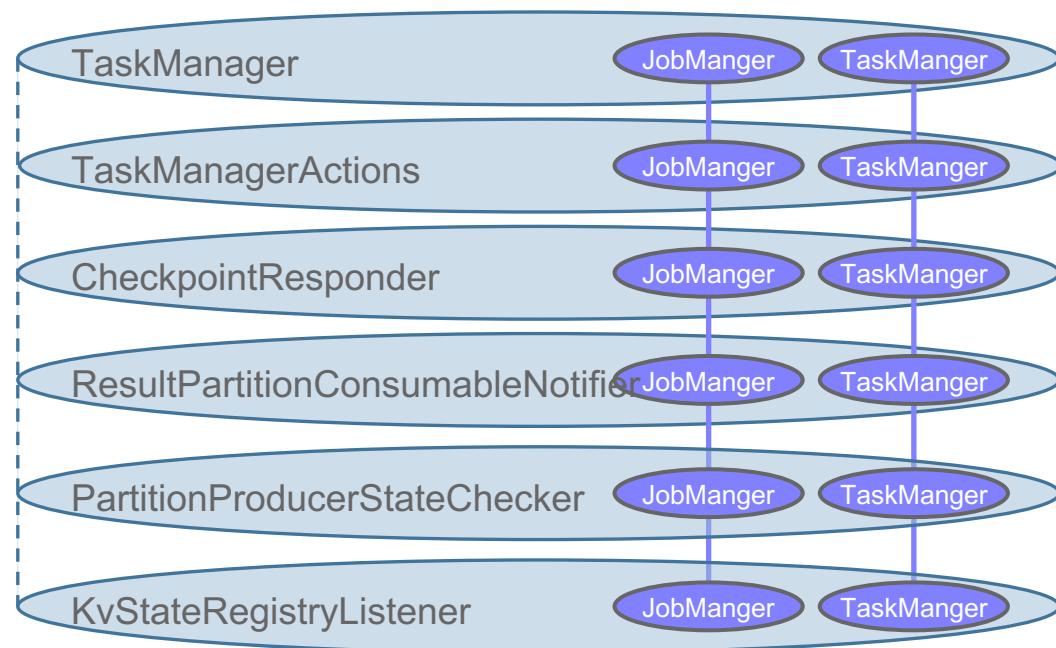
```
@multitier trait Offloading[T] {
  @peer type Master <: { type Tie <: Multiple[Worker] }
  @peer type Worker <: { type Tie <: Single[Master] }
  def run(task: Task[T]): Future[T] on Master =
    placed { (remote(selectWorker())) call execute(task)).asLocal }
  private def execute(task: Task[T]): T on Worker =
    placed { task.process() }
}

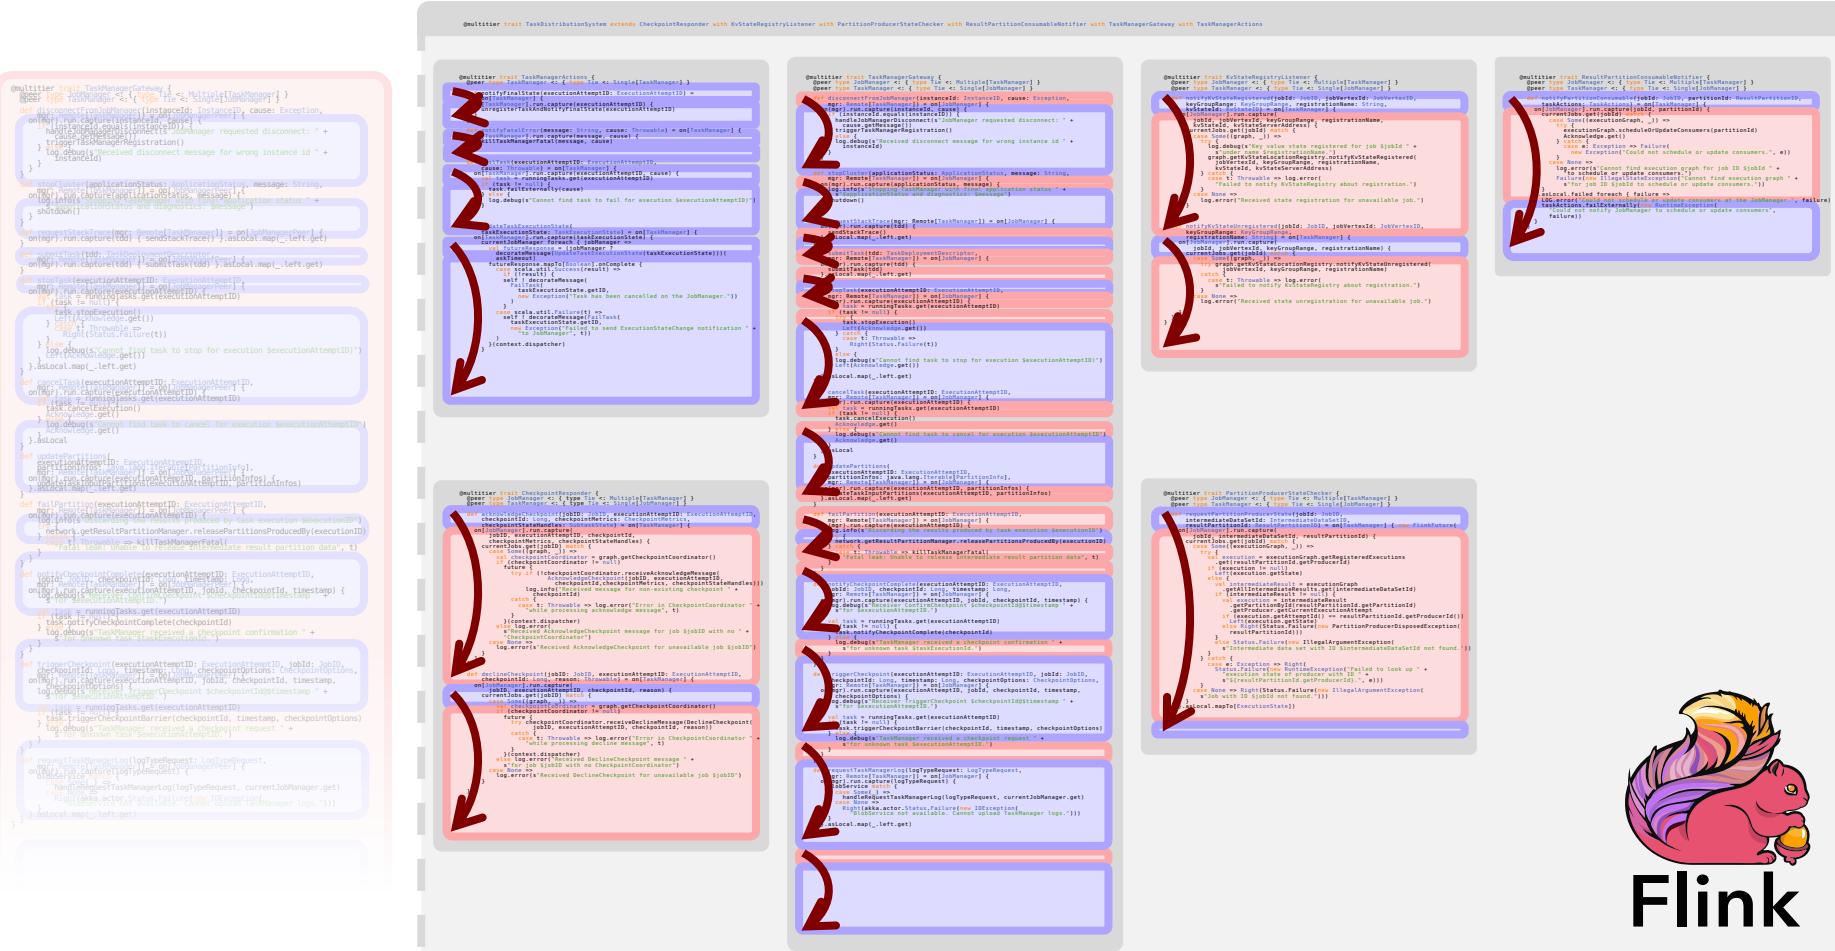
@multitier trait Monitoring {
  @peer type Master <: { type Tie <: Multiple[Worker] }
  @peer type Worker <: { type Tie <: Single[Master] }
  def monitoredTimedOut(monitored: Remote[Worker]): Unit on Master
}

@multitier trait TaskScheduler[T] extends
  Offloading[T] with
  Monitoring
```

Flink Case Study

```
@multitier object TaskDistributionSystem extends
  TaskManager with
  TaskManagerActions with
  CheckpointResponder with
  ResultPartitionConsumableNotifier with
  PartitionProducerSta
  KvStateRegistryListe
```





PRIVACY

The Case of Cambridge Analytica



Local World Business Sports Style/Life Opinion Property Education Media & Marketing

Cambridge Analytica admits influencing polls in Malaysia

Posted on 20 March 2018 - 10:13am

Last updated on 20 March 2018 - 12:42pm

SHARE



Zuckerberg admits his own Facebook data was given to Cambridge Analytica

Cambridge Analytica, accused of enabling Facebook data to be used to manipulate voting intentions in the US and around the world. put to him yesterday: "Are you willing to change your business model to protect users' privacy?"

Zuckerberg, 33, gave one of many evasive responses: "Congressmen, he had the



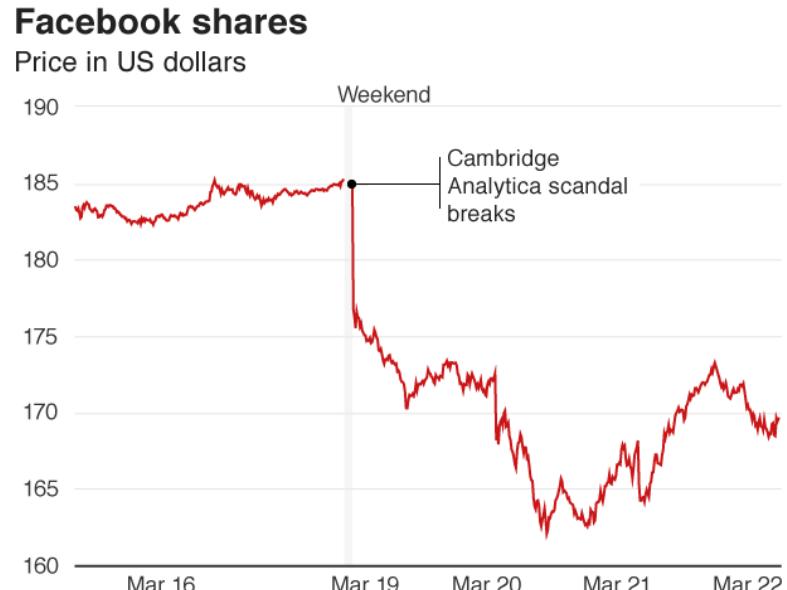
have been on a mission since the Cambridge Analytica scandal, which was brokered by the O'Conor's sister, Sundi. The app called Thinkbait, created by a couple from his wife's University, Cambridge, worked with Cambridge University's centre on ways to use Congress and if Facebook could be held

The Case of Cambridge Analytica

87 Million users affected (FB estimation)

100 billion loss, FB market cap.

Zuckerberg testified in front of the US Congress.



Source: Bloomberg. Last update: 11:20 - 22/03/2018

BBC

Turns out Cambridge Analytica assisted with Trump and Brexit campaign



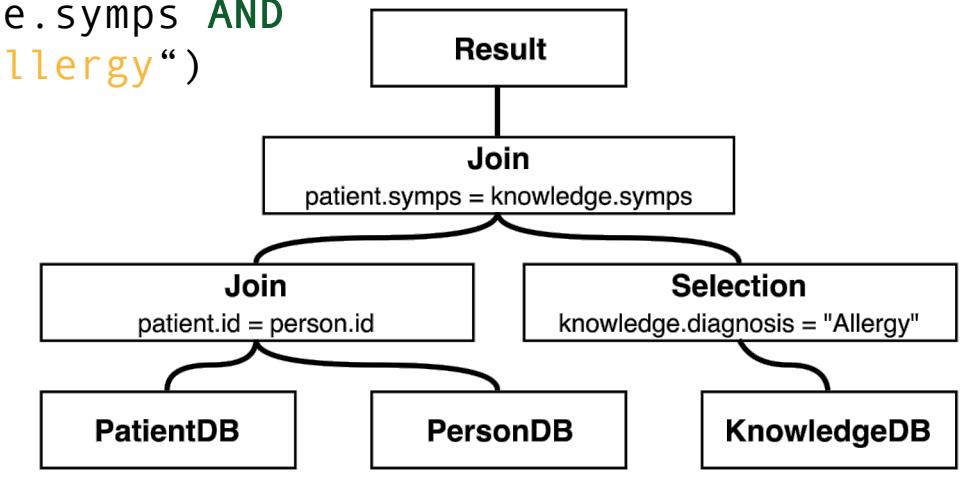
Lesson Learned

We need to engineer software
to ***precisely*** control
where (private) data go!



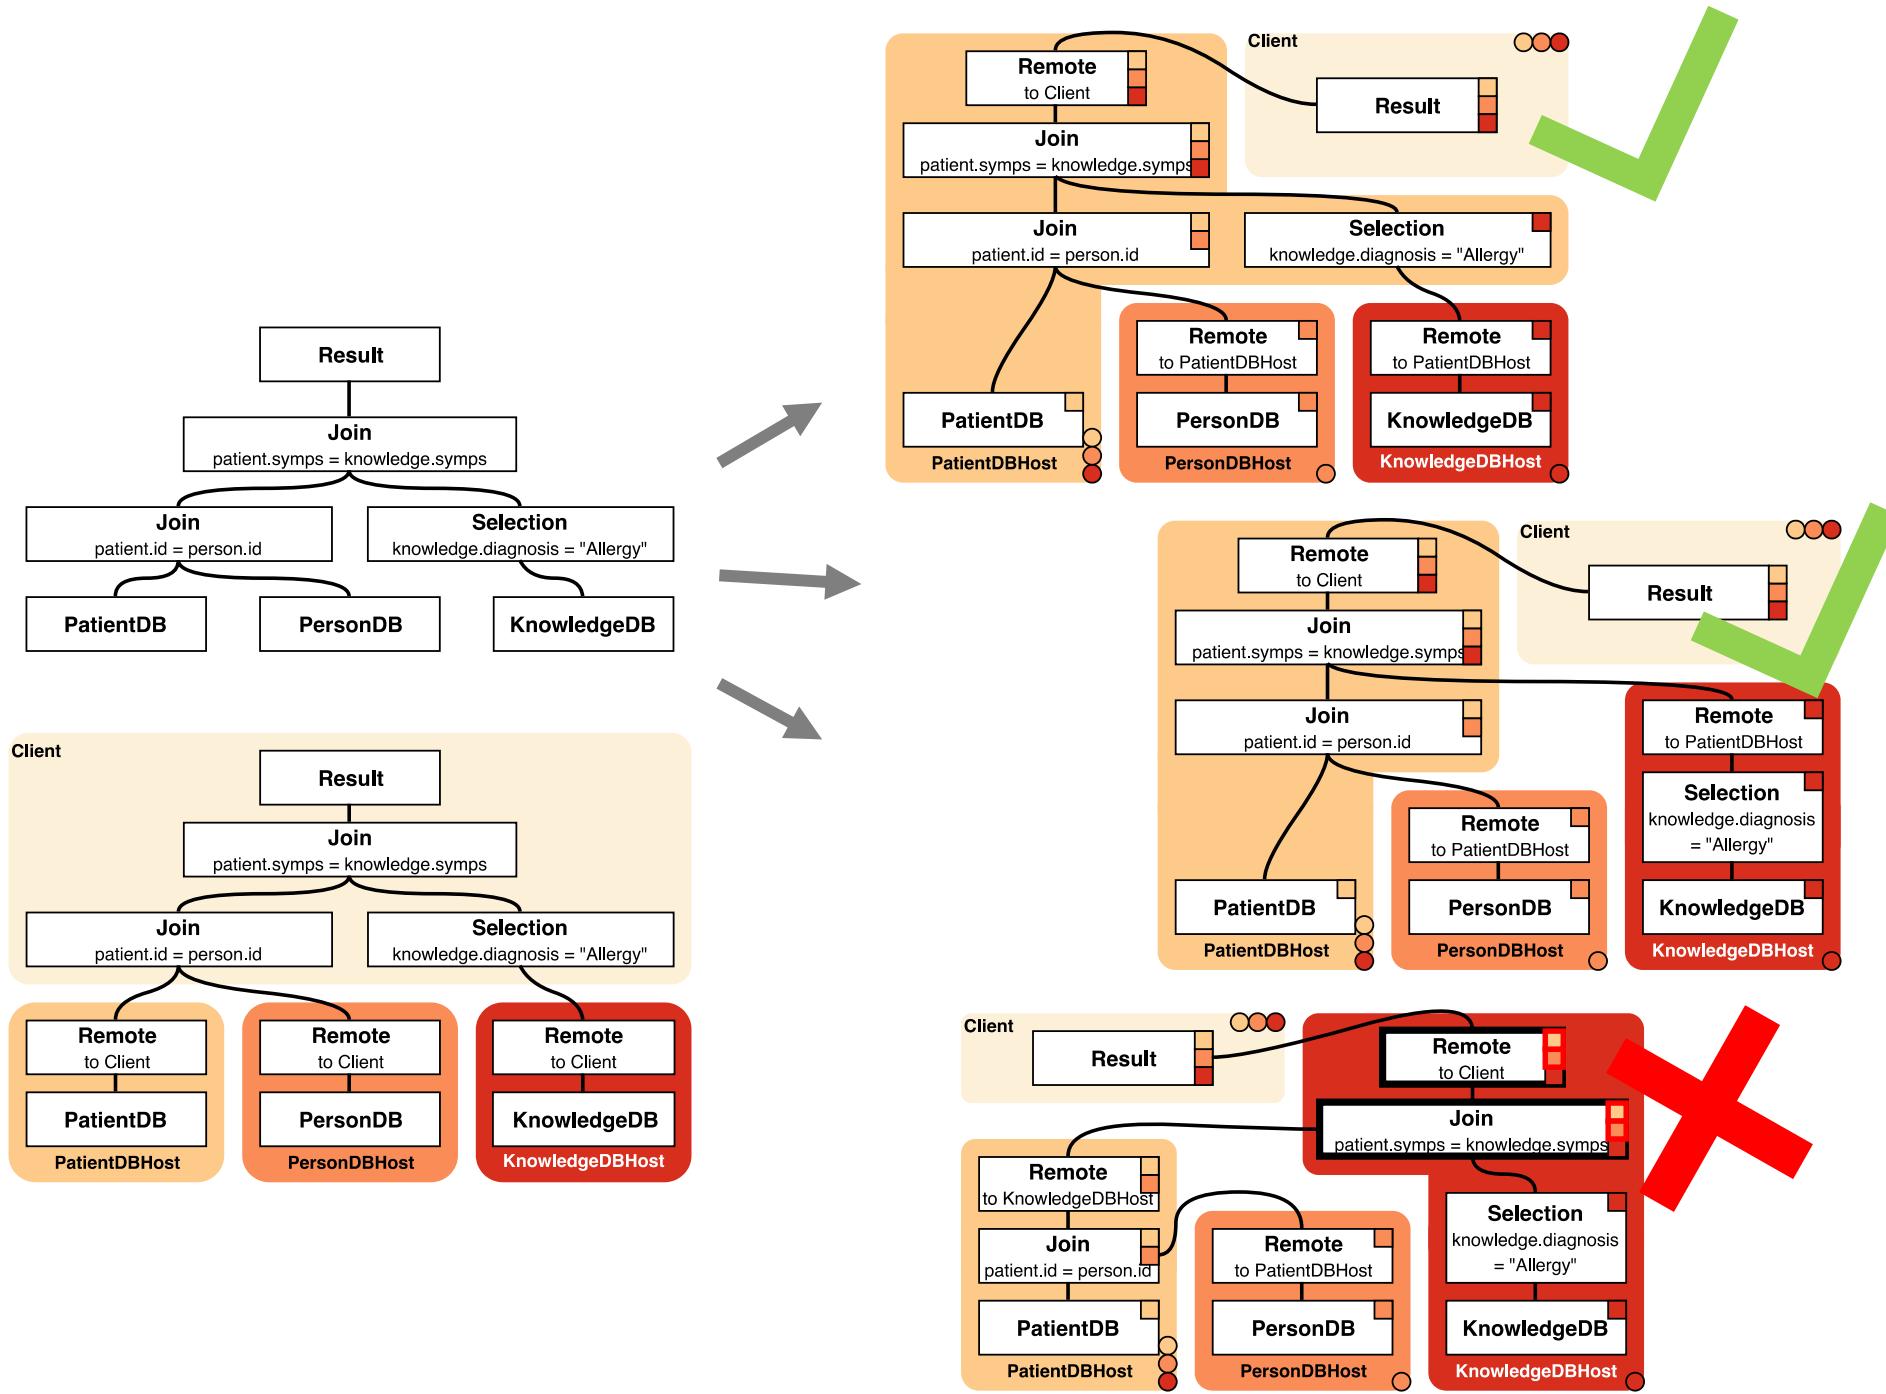
SecQL: Privacy Preserving Queries

```
val result = SELECT (*)
  FROM (personDB,patientDB,knowledgeDB)
 WHERE ((person,patient,knowledge) =>
  person.id == patient.id AND
  patient.symps == knowledge.symps AND
  knowledge.diagnosis == "Allergy")
```



[G.Salvanesci, M.Köhler, D.Sokolowski, P.Haller, S.Erdweg, M.Mezini, **Language-Integrated Privacy-Aware Distributed Queries**, OOPSLA'19]

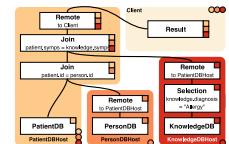
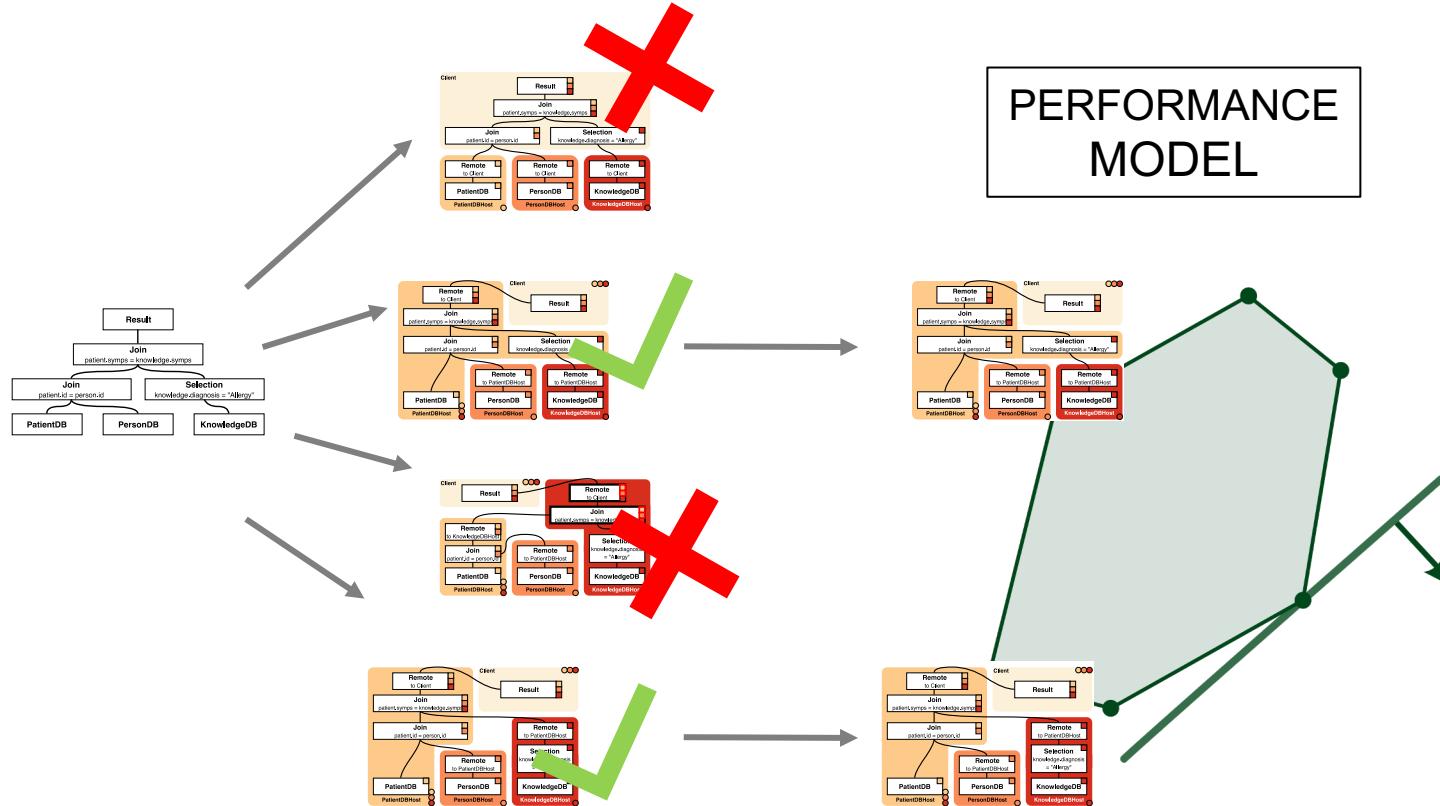
[Ralf Mitschke, Sebastian Erdweg, Mira Mezini, Mirko Kohler, Guido Salvanesci, **i3QL: Language-Integrated Live Data Views**, OOPSLA'14.]



SecQL: 2 Steps Placement Strategy

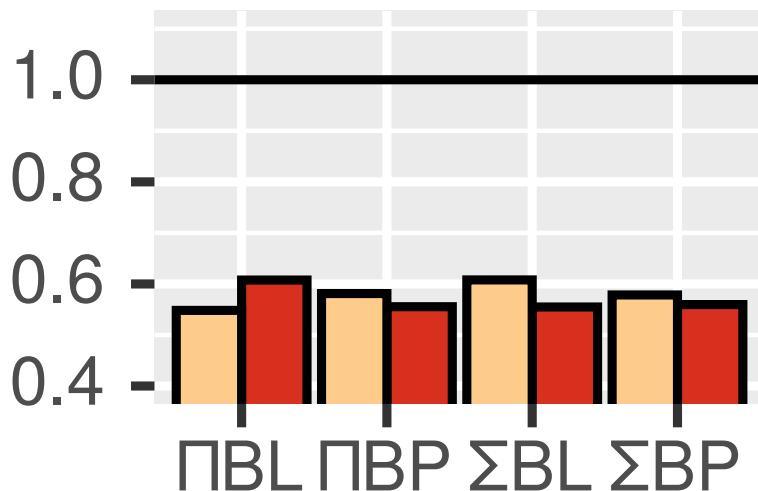
Placement Space Reduction

Cost Model Optimization

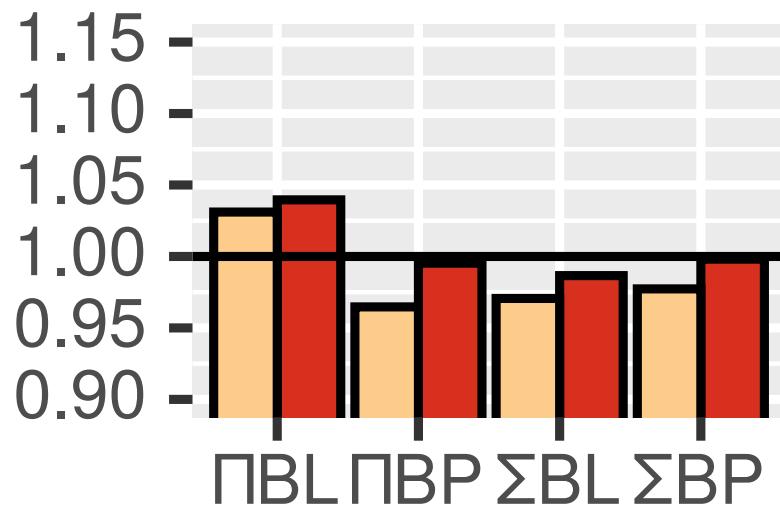


Secure Optimal Placement

Run Time



Memory Growth



■ non-privacy preserving placement ■ privacy preserving placement — client-only baseline

Privacy preserving placement has performance comparable to the optimal placement

FAULT TOLERANCE



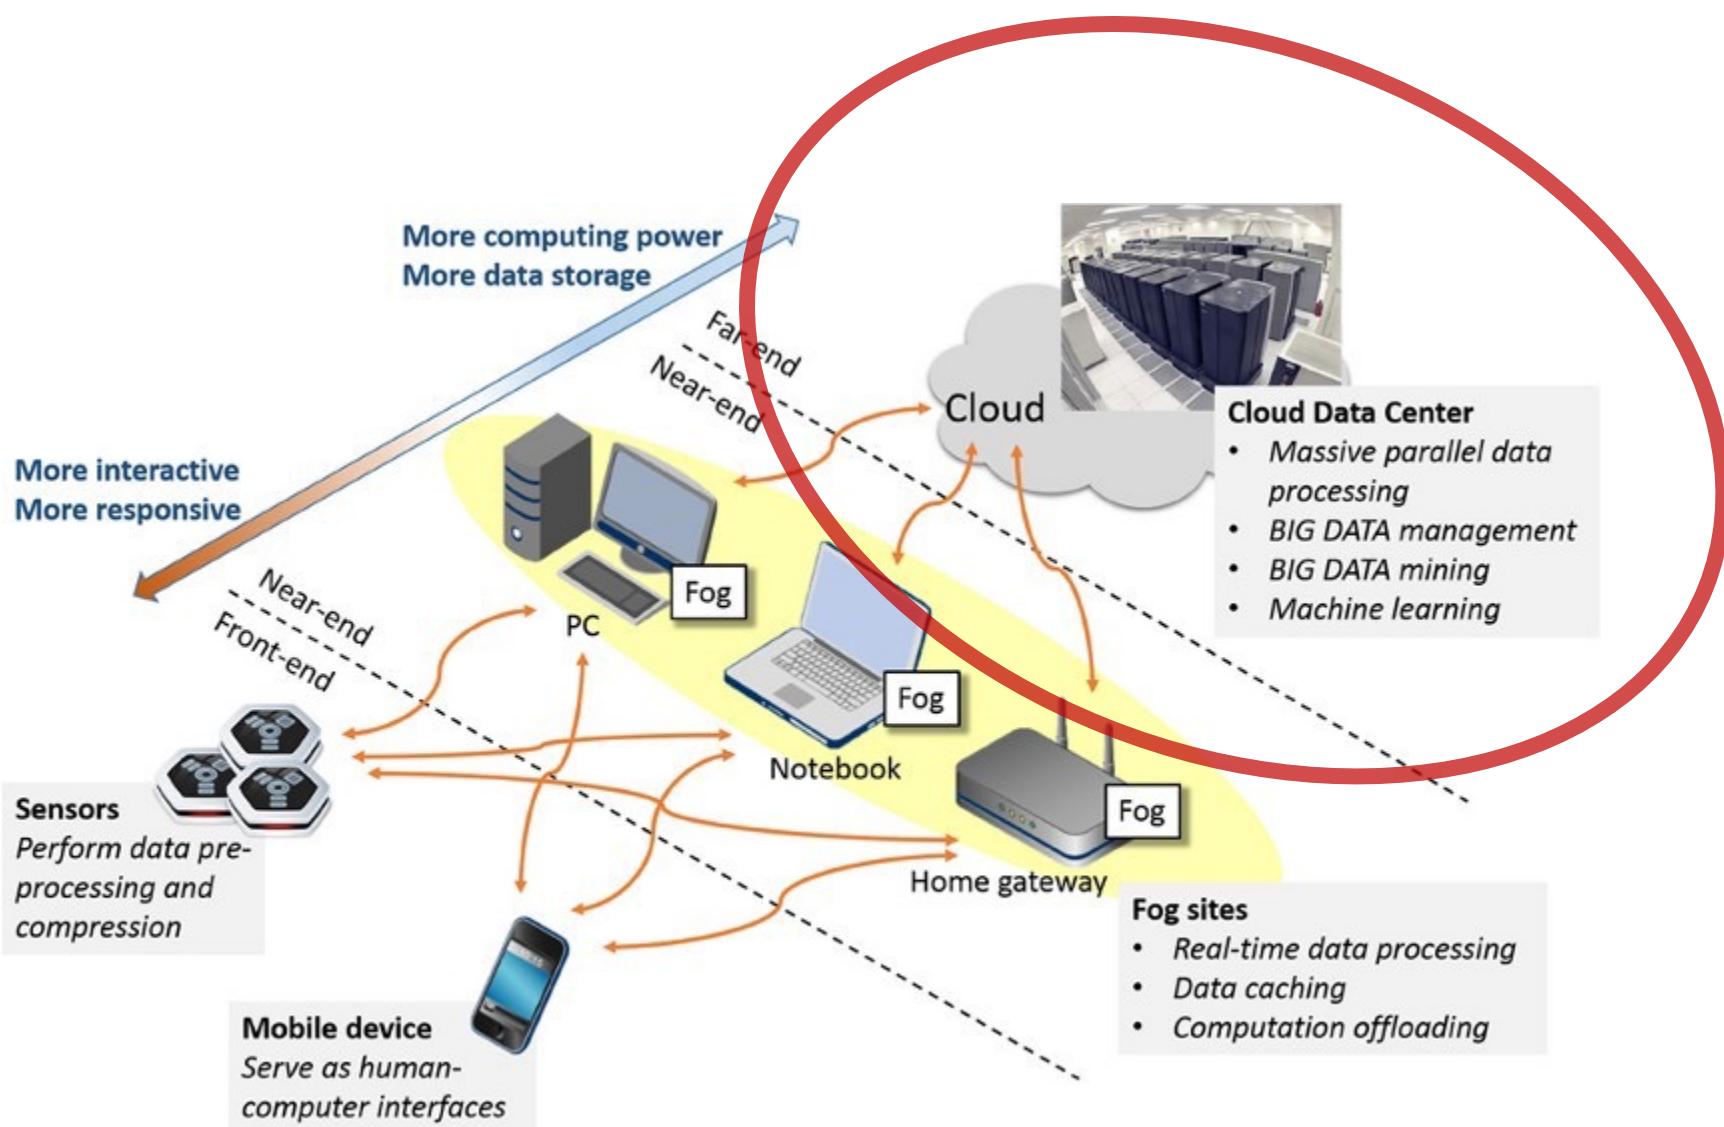
Faults at Facebook

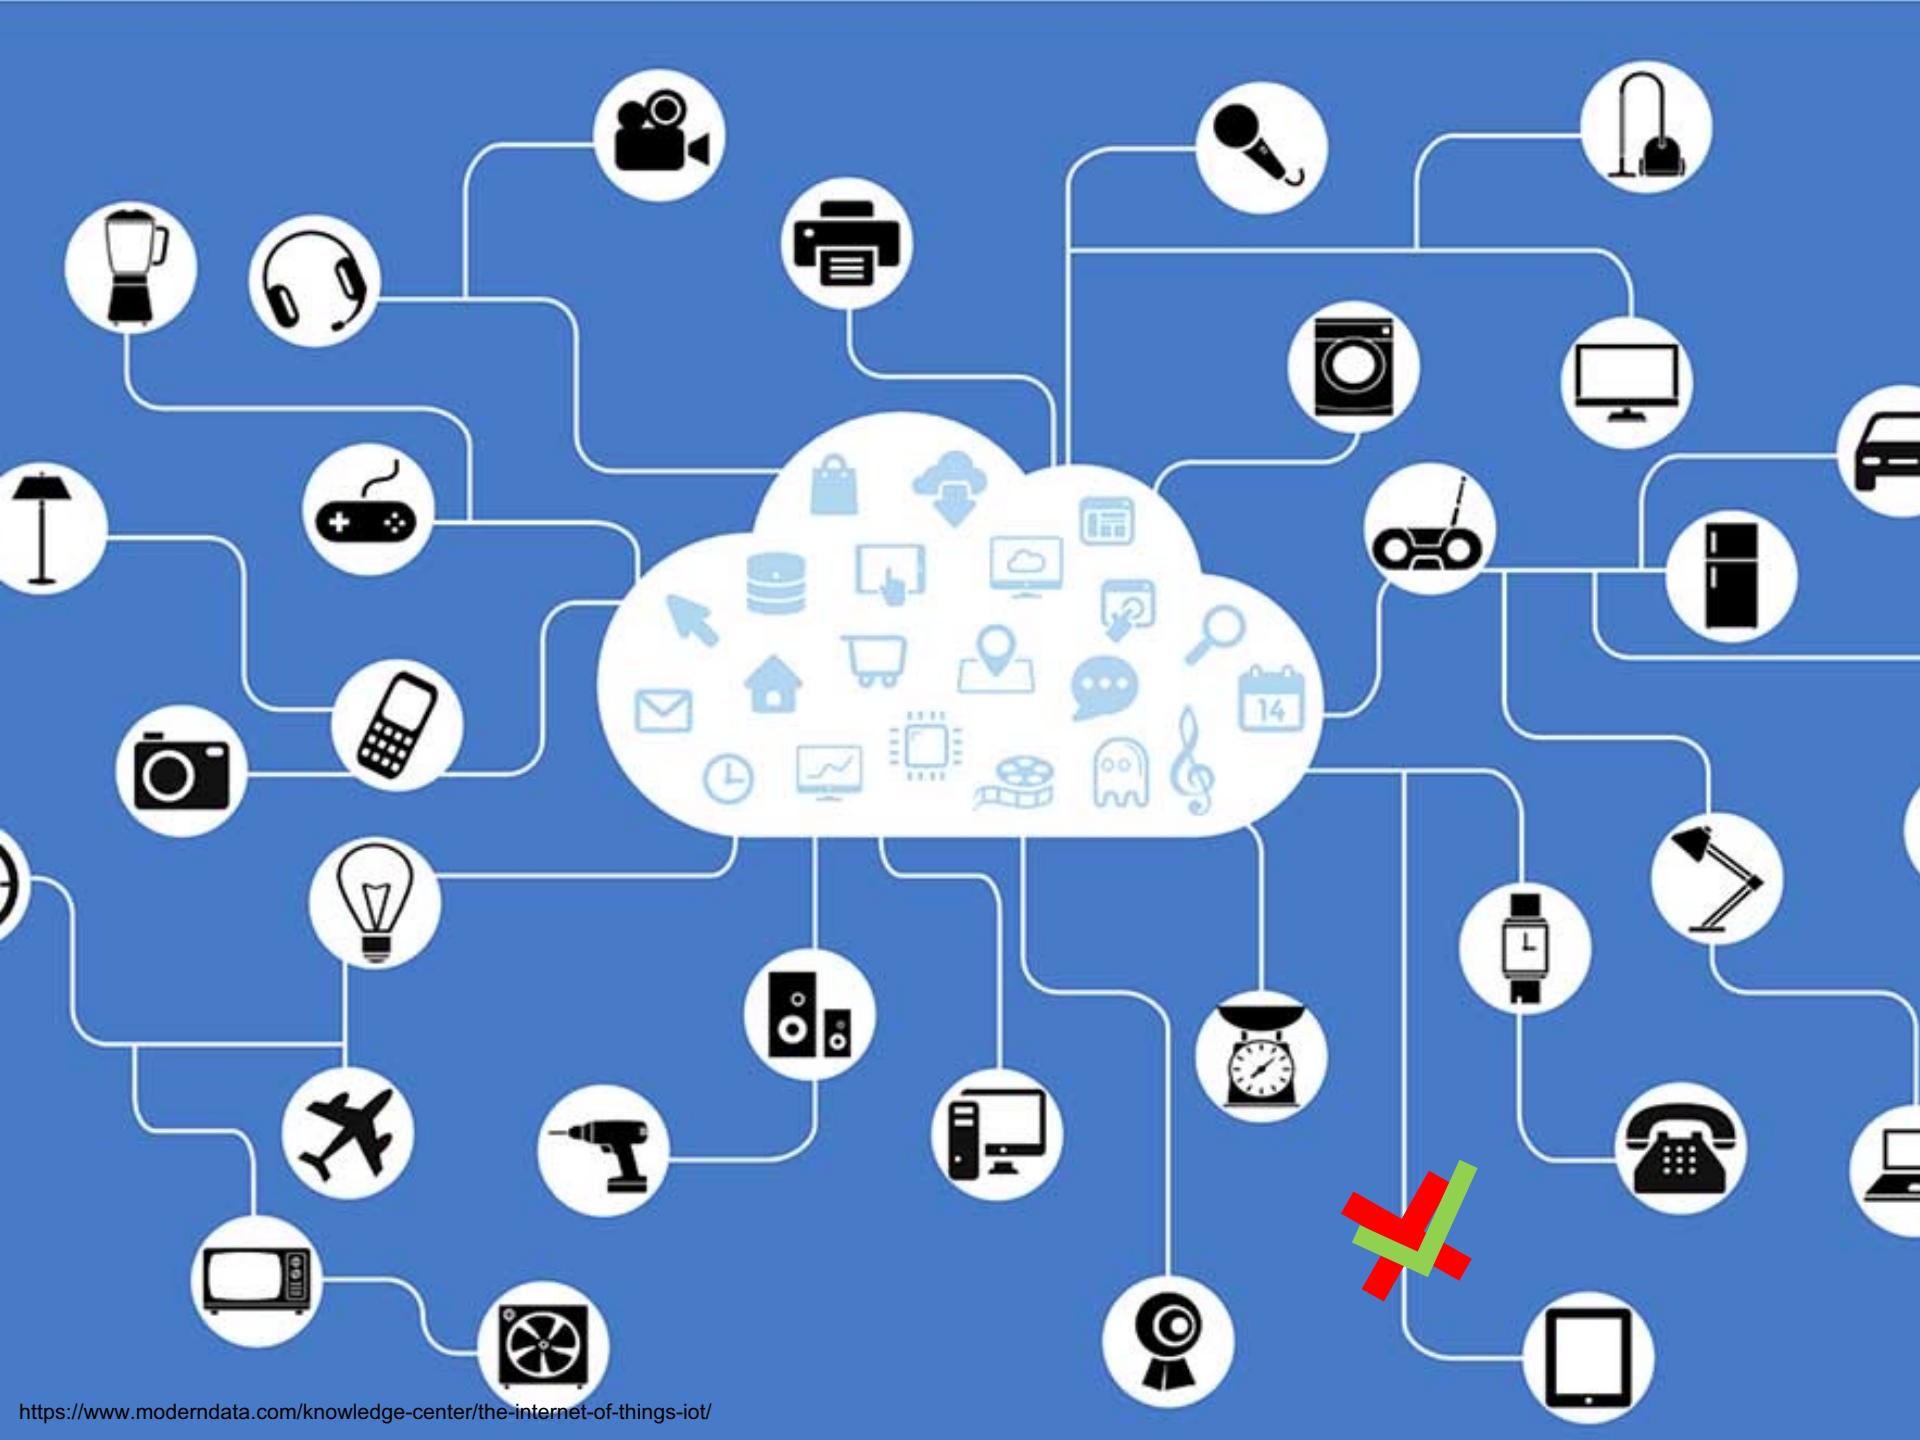


15+ datacenters

100+ data centers outages in 4 years

Planned outages, e.g., Hurricane Florence (Sept'18)





Programming for Failures in Edge Computing

Functional Reactive Animation
Conal Elliott
Microsoft Research
Graphics Group
onal@microsoft.com

Paul Hudak
Yale University
Dept. of Computer Science
paul.hudak@yale.edu

(ation) is a collection of data in Fran are its notions of are time-varying, reactive arbitrarily complex conditions. Most traditional forms, and when images are ions. Although these no other than a programming denotational semantics, at time, to guide reasoning to effectively and effing interval analysis is arital information struc Fran has been implemented for good performance for examples are given, in al phenomena involving

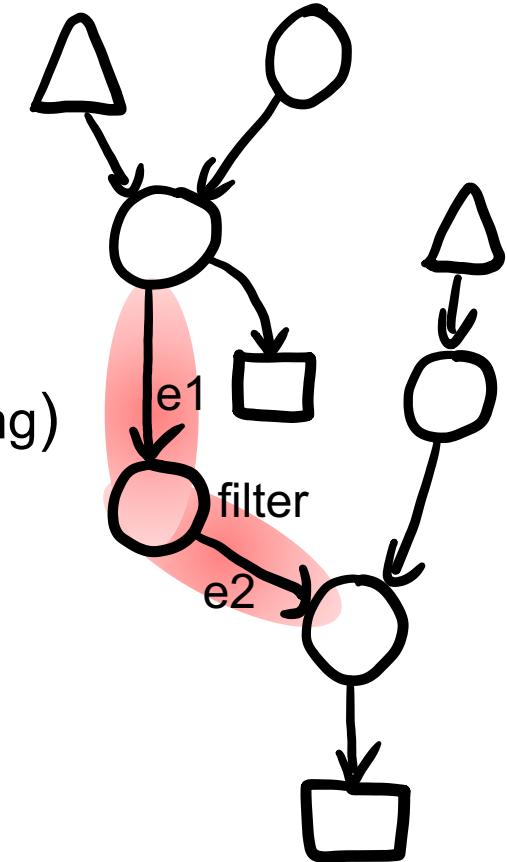
- capturing and handling sequences of motion input events, even though motion input is conceptually continuous;
- time slicing to update each time-varying animation parameter, even though these parameters conceptually vary in parallel; and

By allowing programmers to express the “what” of an “how” of its presentation. With this point of view, it should not be surprising that a set of richly expressive recursive data types, combined with a declarative programming language, serves comfortably for modeling animations, in contrast with the common practice of using imperative languages to program in the conventional hybrid modeling/presentation style. Moreover, we have found that non-strict semantics, higher-order functions, strong polymorphic typing, and systematic overloading are valuable language properties for supporting modeled animations. For these reasons, Fran provides these data types in the programming language, *Functional Fran*.

$e2 = e1.filter(x > 10)$

$e3 = e2.map(x.toString)$

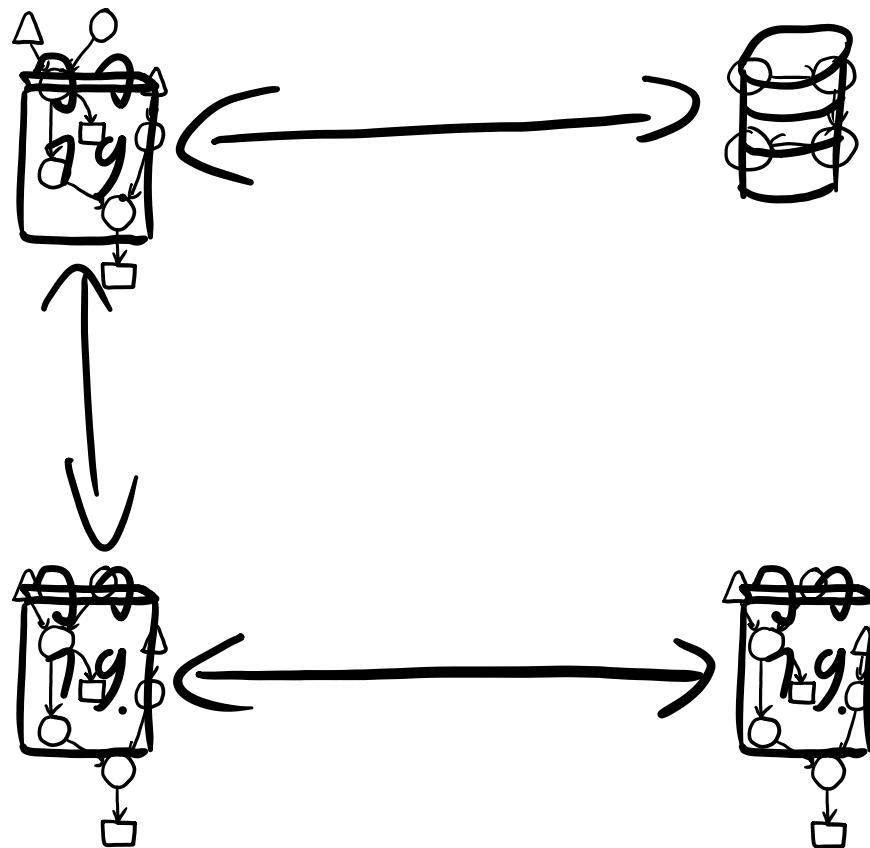
fire $e1(10)$



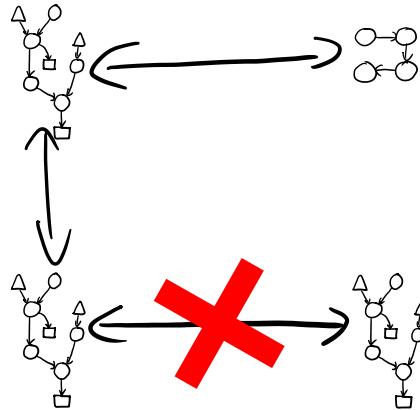
[R.Mogk, J.Drechler, G.Salvanesci, M.Mezini, **A Fault-tolerant Programming Model for Distributed Interactive Applications**, OOPSLA 2019]

[R.Mogk, L.Baumgaertner, G.Salvanesci, B.Freisleben, M.Mezini, **Fault-tolerant distributed reactive Programming**, ECOOP 2018]

Distributed Graphs

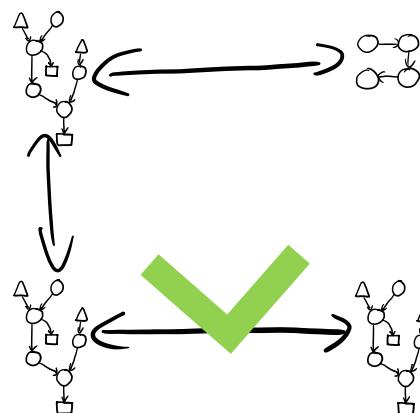


Fault Tolerance „for Free“!



Automatic restore after a crash.

- **Incremental** snapshots of stateful nodes
- Stateless nodes simply ‘**replayed**’



What if your code
is not in reactive style?

Automated Refactoring to Reactive Programming

Mirko Köhler

Reactive Programming Technology
Technische Universität Darmstadt
Darmstadt, Germany
koehler@cs.tu-darmstadt.de

Guido Salvaneschi

Reactive Programming Technology
Technische Universität Darmstadt
Darmstadt, Germany
salvaneschi@cs.tu-darmstadt.de

Abstract—

Reactive programming languages and libraries, such as ReactiveX, have been shown to significantly improve software design and have seen important industrial adoption over the last years. Asynchronous applications – which are notoriously error-prone to implement and to maintain – greatly benefit from reactive programming because they can be defined in a declarative style, which improves code clarity and extensibility.

In this paper, we tackle the problem of refactoring existing software that has been designed with traditional abstractions for asynchronous programming. We propose 2Rx, a refactoring approach to automatically convert asynchronous code to reactive programming. Our evaluation on top-starred GitHub projects shows that 2Rx is effective with common asynchronous constructs and it can provide a refactoring for 91.7% of their occurrences.

Keywords-refactoring; asynchronous programming; reactive programming; Java;

over low level abstractions like threads, but come with their own limitations. For example, AsyncTask does not easily support composition, like sequencing multiple asynchronous computations.

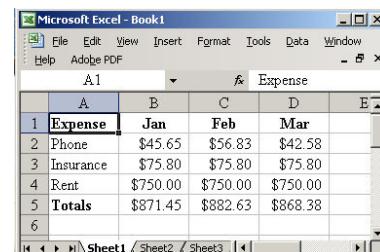
Recently, Reactive Programming (RP) has emerged as a programming paradigm specifically addressing software that combines events [3]. Crucially, RP allows to easily express computations on event streams that can be chained and combined using high-order functional operators. This way, each operator can be scheduled independently, providing a convenient model for asynchronous programming. As a result, RP provides means to describe asynchronous programs in a declarative way. Previous research showed that RP can support the software design of reactive systems [1, 2].

REACTIVE PROGRAMMING

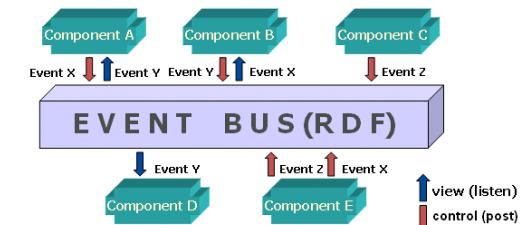
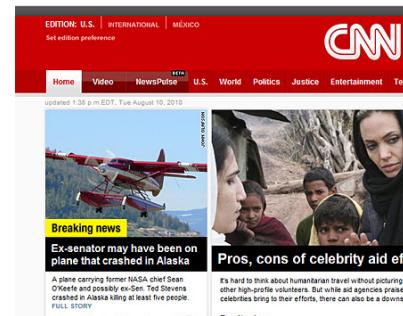
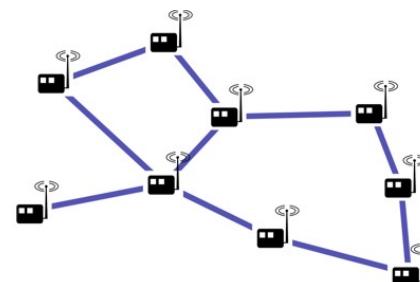
Reactive Software is Ubiquitous

Web apps, IDEs, monitoring systems, ...

- Detect events
- Combine reactions
- Propagate changes
- Aggregate/filter



A	B	C	D	E
1	Expense	Jan	Feb	Mar
2	Phone	\$45.65	\$56.83	\$42.58
3	Insurance	\$75.80	\$75.80	\$75.80
4	Rent	\$750.00	\$750.00	\$750.00
5	Totals	\$871.45	\$882.63	\$868.38
6				



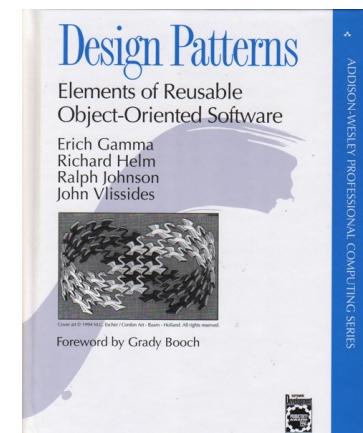
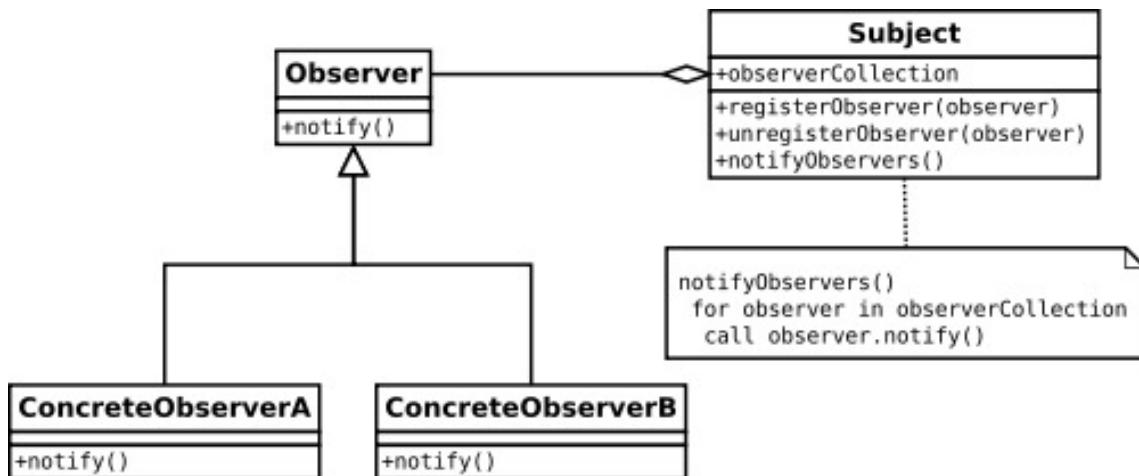
Reactive Applications are not Easy to Develop



Event handling:

- 30% of code in desktop applications
- 50% of bugs reported during production cycle

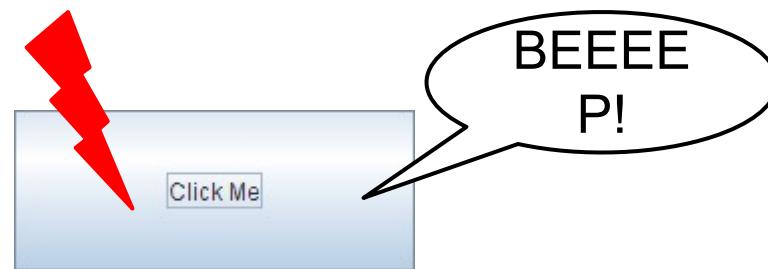
The (good? old) Observer Pattern



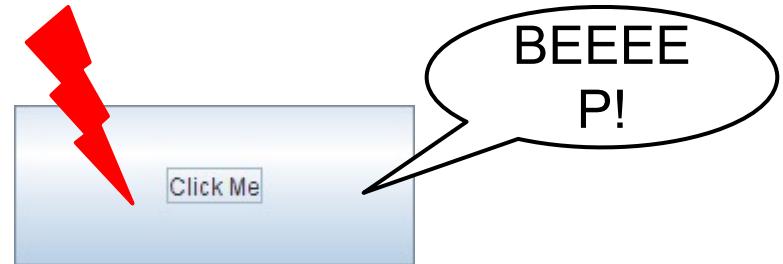
The Observer Pattern

What about Java Swing ?

- javax.swing



```
public class Beeper extends JPanel implements ActionListener {  
    JButton button;  
  
    public Beeper() {  
        super(new BorderLayout());  
        button = new JButton("Click Me");  
        button.setPreferredSize(new Dimension(200, 80));  
        add(button, BorderLayout.CENTER);  
        button.addActionListener(this);  
    }  
  
    public void actionPerformed(ActionEvent e) {  
        Toolkit.getDefaultToolkit().beep();  
    }  
  
    private static void createAndShowGUI() { // Create the GUI and show it.  
        JFrame frame = new JFrame("Beeper"); //Create and set up the window.  
        frame.setDefaultCloseOperation(JFrame.EXIT_ON_CLOSE);  
        JComponent newContentPane = new Beeper(); //Create and set up the content pane.  
        newContentPane.setOpaque(true);  
        frame.setContentPane(newContentPane);  
        frame.pack(); //Display the window.  
        frame.setVisible(true);  
    }  
  
    public static void main(String[] args) {  
        javax.swing.SwingUtilities.invokeLater( new Runnable() { public void run() {createAndShowGUI();}});  
    }  
}
```



Reactive Programming

Can **language abstractions** help software design?

Events

Signals

Similar examples:

fault-tolerance -> **exceptions**

encapsulation -> visibility **modifiers**

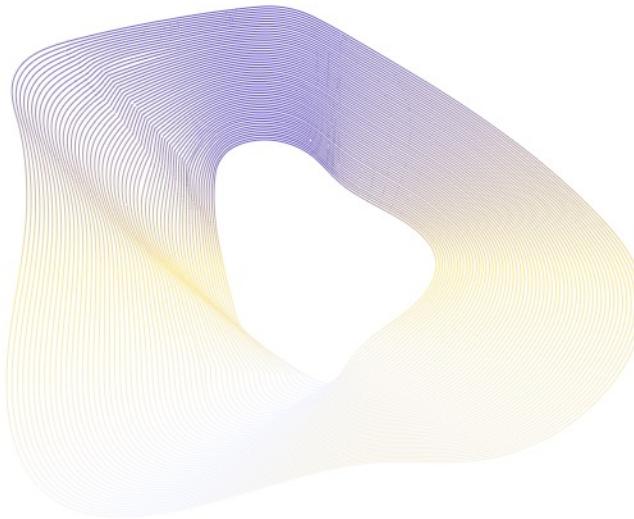
extensibility -> **inheritance**



Reactive programming

REScala is the most advanced solution for functional reactive programming on the JVM and the Web. It provides you with both, the latest features directly from academic research with a stable well maintained core.

REScala's features include transactional guarantees, eventual consistency, event and signal APIs, reactive HTML DOM interactions, compatibility with reactive streams, and many more.

[Get Started](#)[More →](#)

www.rescala-lang.com



For Academics

REScala offers the latest and greatest from research on functional reactive programming embedded into an object-oriented language combined with current research on distributed languages.



For Industry

REScala has a stable API and [regular releases](#). We have a large test suite and many internal case studies. Source incompatible changes are rare and only happen to improve ergonomics for very common cases.



For Students

REScala is perfect for learning reactive programming concepts and learning how to design Scala libraries in general. It is also easy to extend, and we offer [theses and other opportunities](#).

Events

Registered handlers are executed when the event fires

- The actual parameter is provided to the handler

```
val e = Evt()  
e += { x => println("The value of the event is: " + x) }
```

```
e(10)
```

```
e(11)
```

-- output ----

The value of the event is: 10

The value of the event is: 11

Constrains as Reactive Values

What about expressing *functional dependencies* as constraints ?

```
val a = 3
val b = 7
val c = a + b // Statement
println(c)
> 10
a= 4
println(c)
> 10
```

```
val a = 3
val b = 7
val c := a + b // Constraint
println(c)
> 10
a= 4
println(c)
> 11
```

Constrains as Reactive Values

Signals: (pure)reactive expressions

Constraints “automatically” enforced

.now returns the current value

Vars: primitive reactive values

Updated “manually”

```
val a = Var(3)  
val b = Var(7)  
val c = Signal{ a() + b()  
}
```

```
println(c.now)  
> 10
```

```
a.set(4)  
println(c.now)  
> 11
```

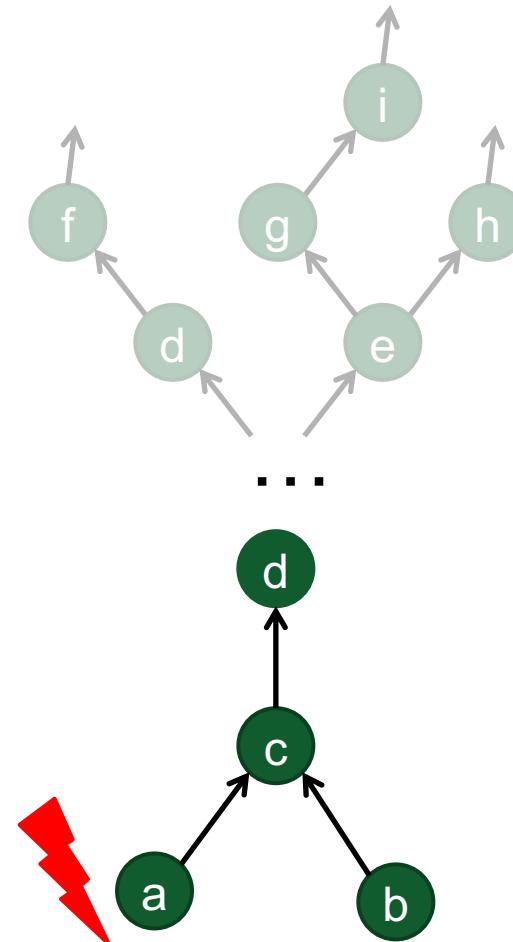
Reference Model

Change propagation model

- Dependency graph
- Push-driven evaluation

```
val a = Var(3)
val b = Var(7)
val c = Signal{ a() + b()
}
val d = Signal { 2 * c() }
```

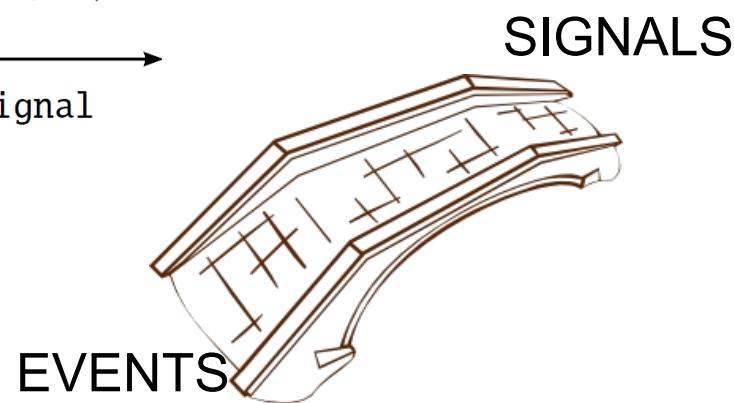
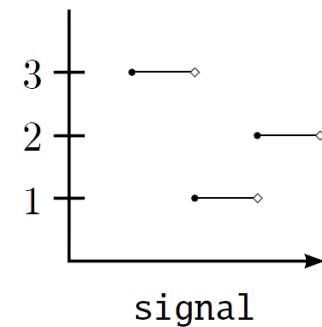
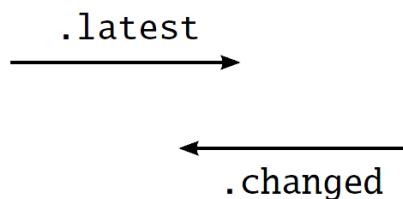
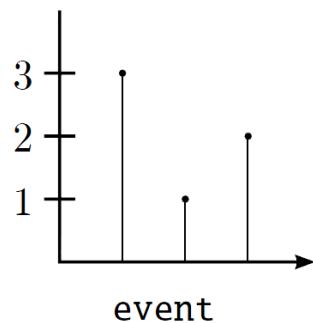
...



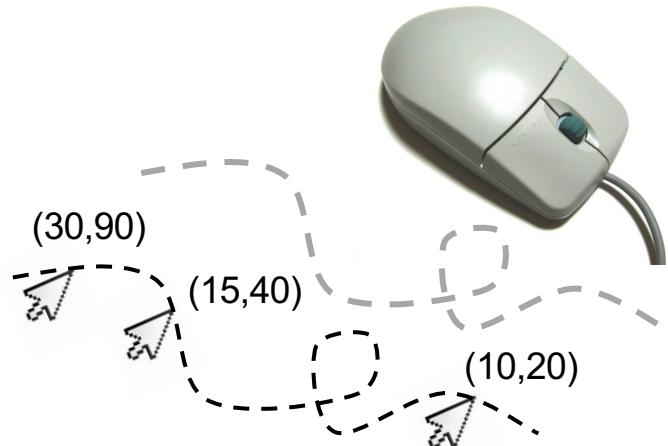
Basic Conversion Functions

`Changed :: Signal -> Event`

`Latest :: Event -> Signal`



Mouse with RP



val position: Signal[(Int,Int)] = mouse.position

val shiftedPosition: Signal[(Int,Int)] = Signal{ mouse.position + (10, 10) }

evt clicked: Event[Unit] = mouse.clicked

val lastClick: Signal[(Int,Int)] = position snapshot clicked

Example: Observer

```
/* Create the graphics */
title = "Reactive Swing App"
val button = new Button {
  text = "Click me!"
}
val label = new Label {
  text = "No button clicks registered"
}
contents = new BoxPanel(Orientation.Vertical) {
  contents += button
  contents += label
}
```



```
/* The logic */
listenTo(button)
var nClicks = 0
reactions += {
  case ButtonClicked(b) =>
    nClicks += 1
    label.text = "Number of button clicks: " +
    nClicks
    if (nClicks > 0)
      button.text = "Click me again"
}
```

Example: Signals

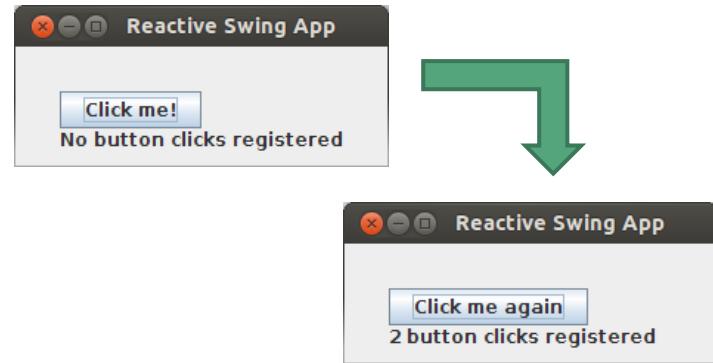
```
title = "Reactive Swing App"  
val label = new ReactiveLabel  
val button = new ReactiveButton
```

```
val nClicks = button.clicked.fold(0) {(x, y) => x + 1}
```

```
label.text = Signal { ( if (nClicks() == 0) "No" else nClicks() ) + " button clicks registered" }
```

```
button.text = Signal { "Click me" + (if (nClicks() == 0) "!" else " again" ) }
```

```
contents = new BoxPanel(Orientation.Vertical) {  
    contents += button  
    contents += label  
}
```



Claim: RP beats OO (Observer)

Easier to compose

Declarative style

Easier program comprehension

State management not explicit

Automatic memory management

Flapjax: A Programming Language for Ajax Applications

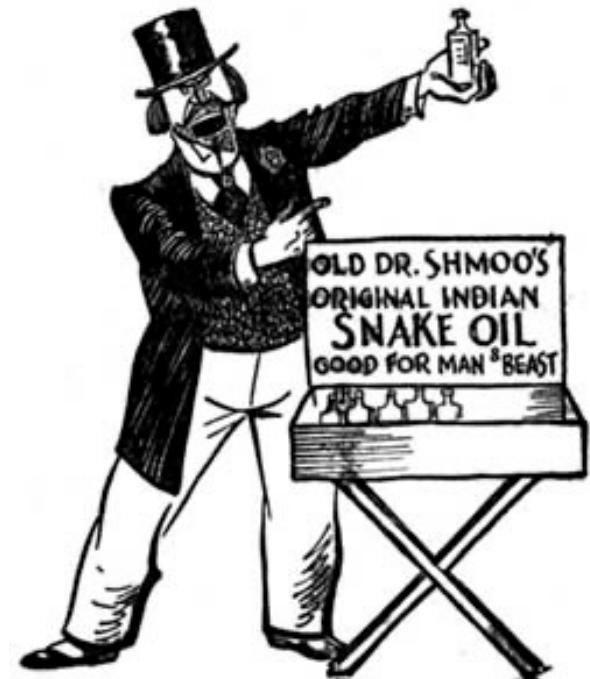
“Obviously, the Flapjax code may not appear any ‘easier’ to a first-time reader”

[Leo A. Meyerovich, Arjun Guha, Jacob Baskin, Gregory H. Cooper, Michael Greenberg, Aleks Bromfield, Shriram Krishnamurthi, *Flapjax: A Programming Language for Ajax Applications, OOPSLA’09*]

Keywords JavaScript, Web Programming, Functional Reactive Programming

1. Introduction

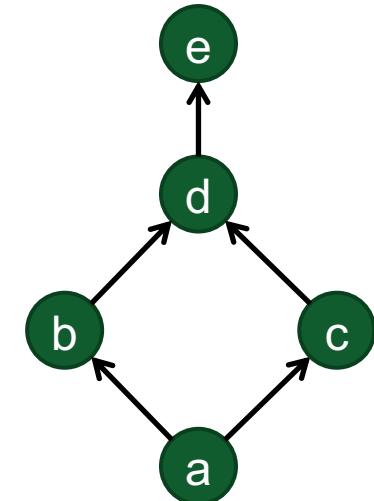
The advent of broadband has changed the structure of application software. Increasingly, desktop applications are migrating to the Web. Programs that once made brief forays



The study

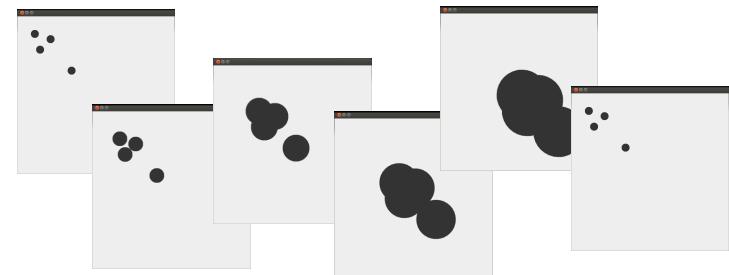
10 applications, ~130 subjects

- RP and OO group (**between** subj.)
 - Questions for comprehension



What to measure?

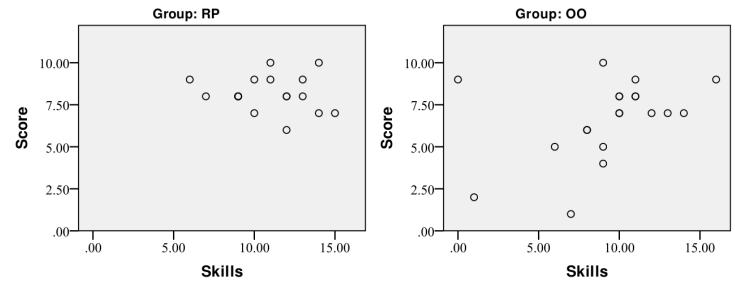
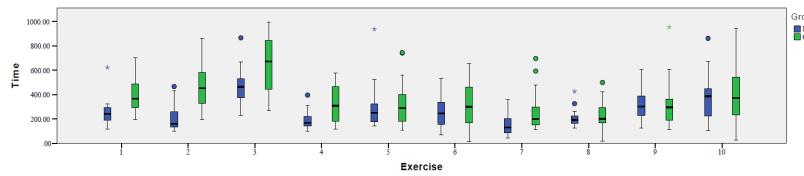
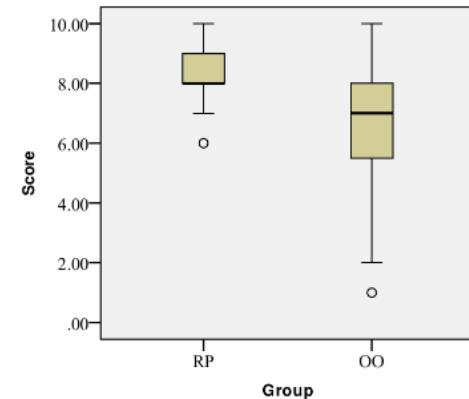
- **Time** to answer a question
 - Amount of **correct answers**



Results

REScala increases correctness of program comprehension

In REScala, comprehension is more time-consuming



[Guido Salvaneschi, Sven Amann, Sebastian Proksch, Mira Mezini, **An Empirical Study on Program Comprehension with Reactive Programming**, FSE'14.]

[G. Salvaneschi, S. Proksch, S. Amann, S. Nadi, M. Mezini, ***On the Positive Effect of Reactive Programming on Software Comprehension: An Empirical Study***, TSE'17]

Teaching Reactive Programming

Master course (9CP)

Software Engineering: Design & Construction

Design patterns

Domain specific languages

Software architectures

Reactive Programming with REScala

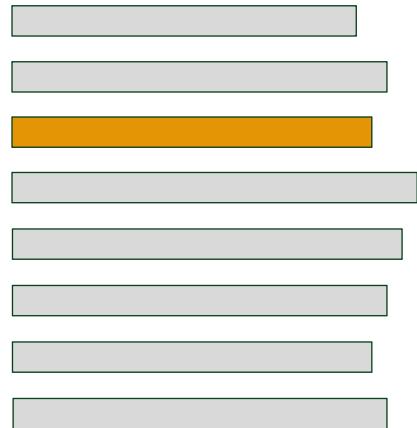
...

HOW TO DEBUG
REACTIVE PROGRAMS !?



Debugging for Reactive Programming

Traditional debugging (Imperative)



```
0x051DE590  a0 e5 1d 05 4c
0x051DE5A0  20 e6 1d 05 b1
0x051DE5B0  40 e6 1d 05 00
0x051DE5C0  bc c3 94 70 b4
```

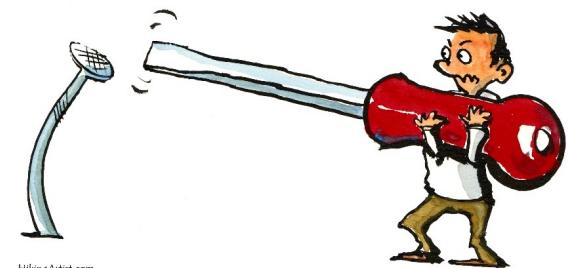
Program Stack

Step over
statements

Inspect state

Reactive Programming (Declarative)

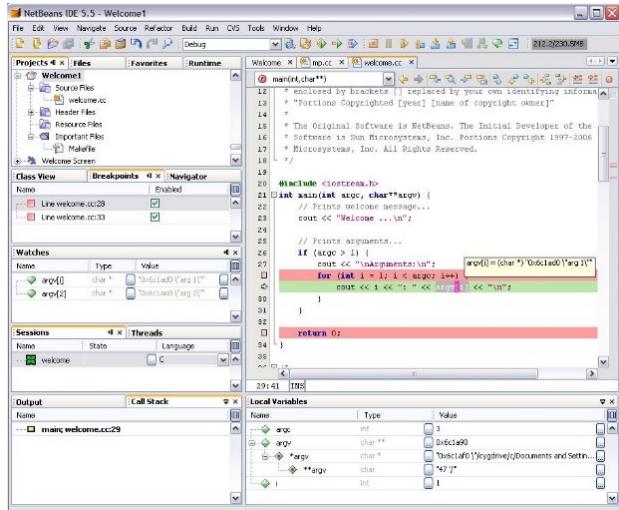
Signals



Abstract
over state

!?

A Paradigm Shift



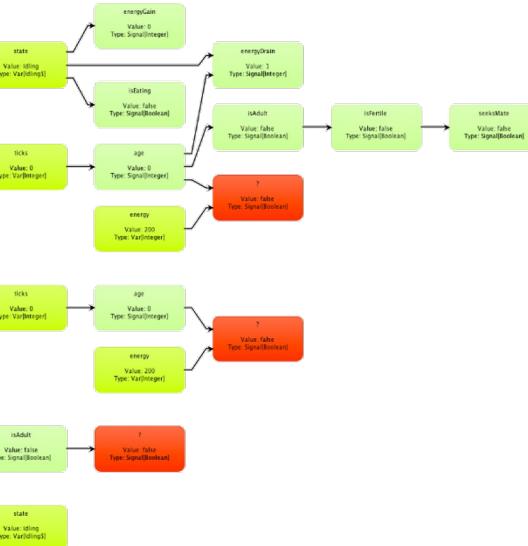
Traditional Debugging

Stepping over statements

Breakpoint on line X

Inspect memory

Navigate object references



RP Debugging

Stepping over the dependency graph

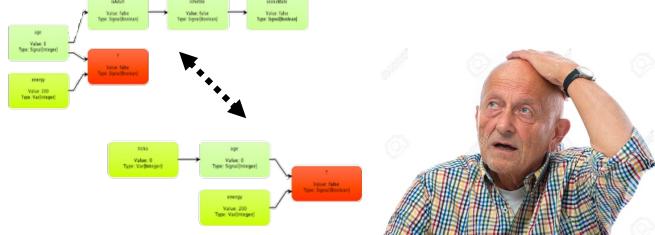
Breakpoint on node X

Inspect values in the dependency graph

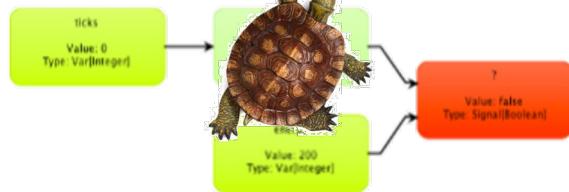
Navigate signals in the graph

Bug Hunting with Reactive Debugging

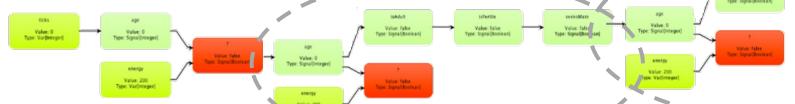
Missing dependencies



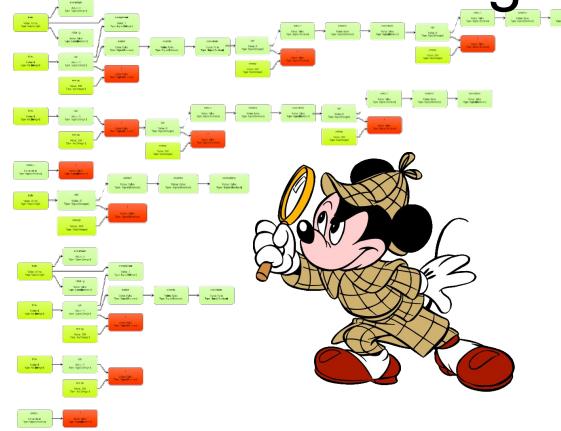
Performance bugs



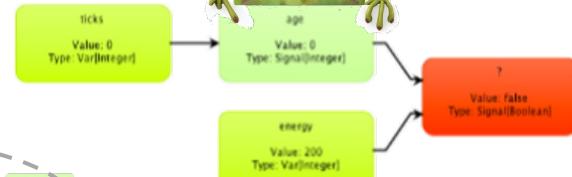
Memory and time leaks



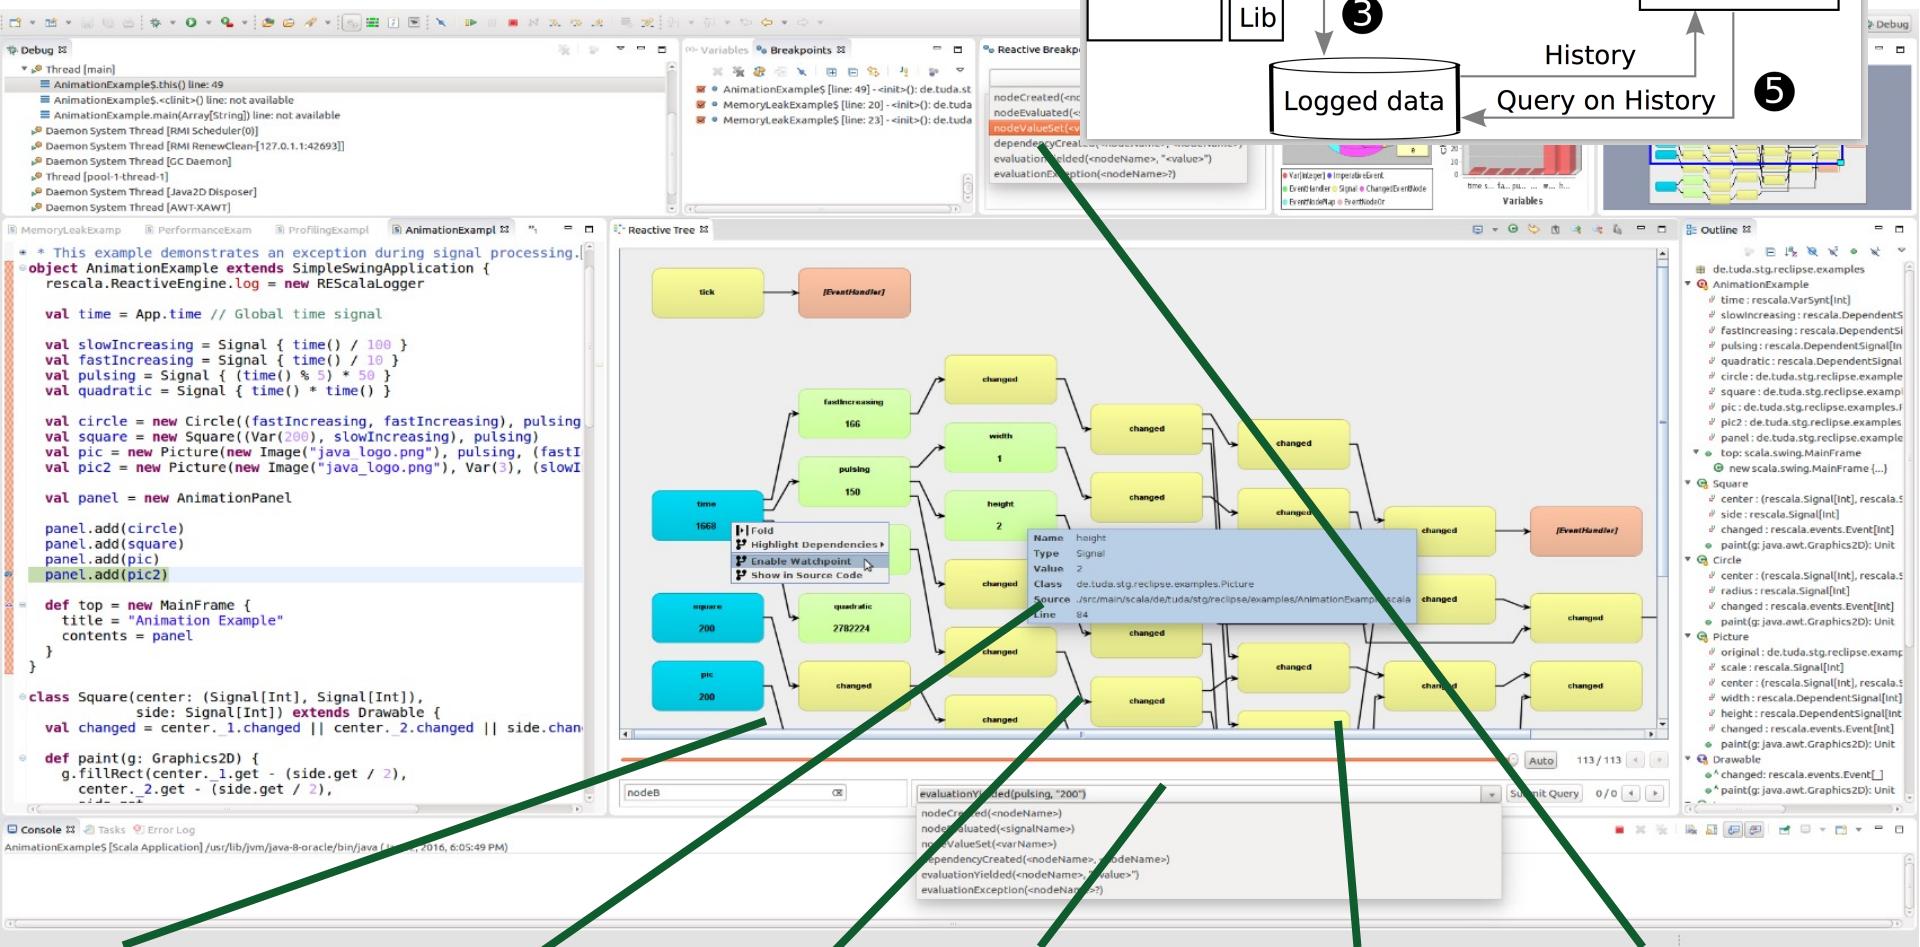
Understanding RP programs



Bugs in signal expressions



Reactive Inspector (Eclipse plugin - Scala IDE)



NODE
SEARCH

NODE
BREAKPOINTS

TREE

NODE
INSPECTION QUERIES

BACK-IN-TIME
REACTIVE
DEBUGGING

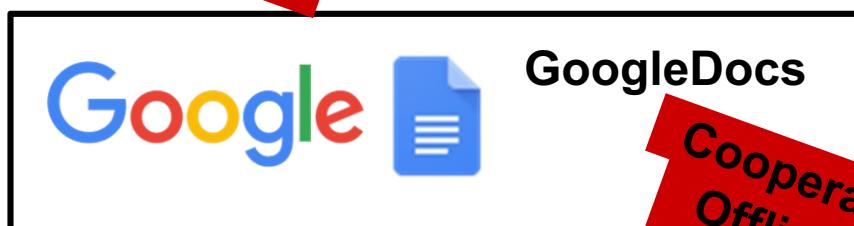
CONSISTENCY

Replicated Data

Replicated data is used to...

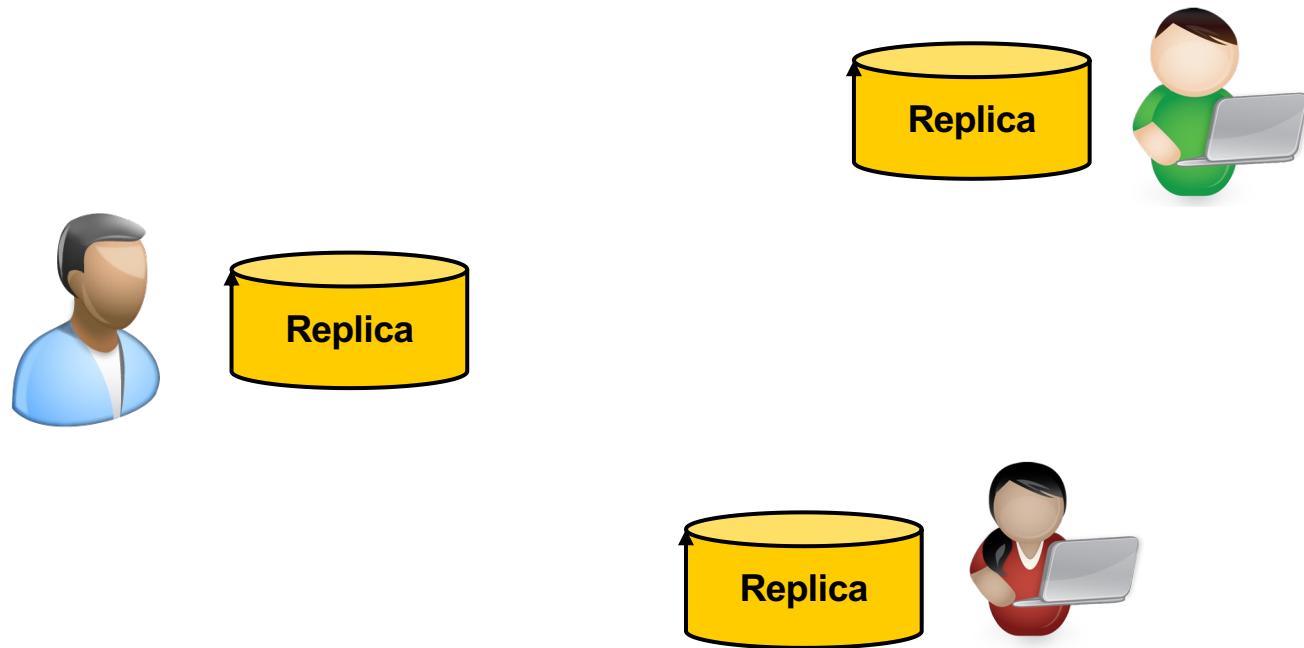


User-experience

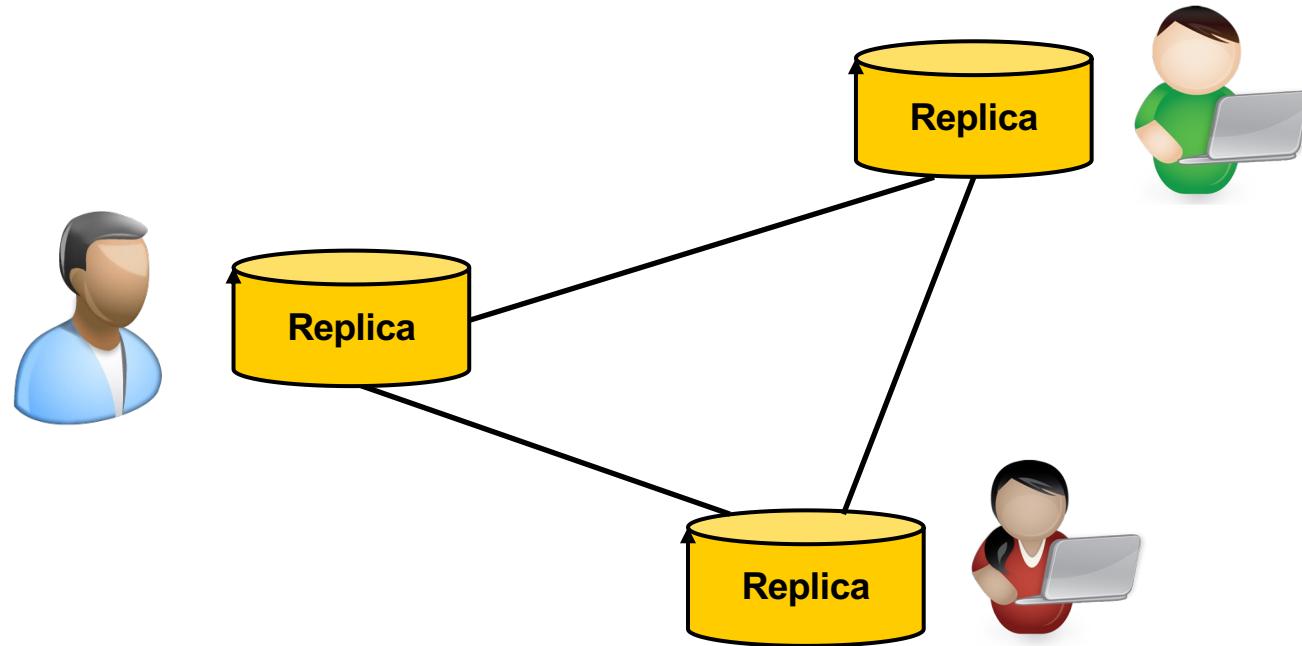


*Cooperative Editing
Offline functionality*

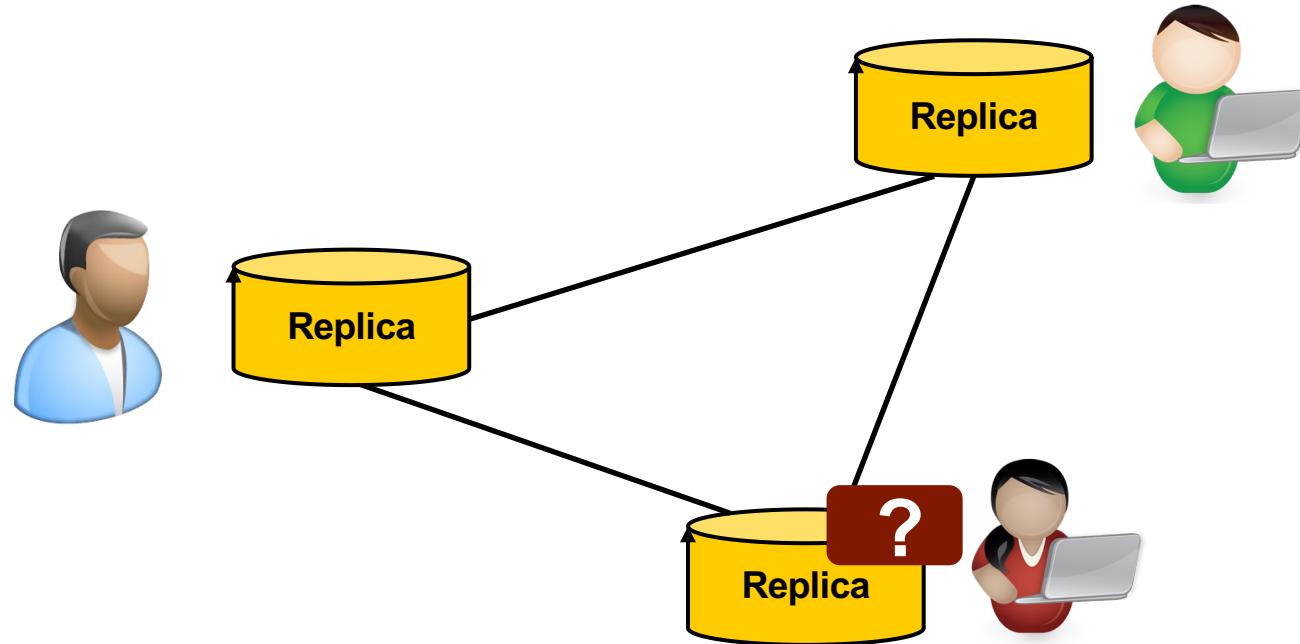
Example: Replicated Data



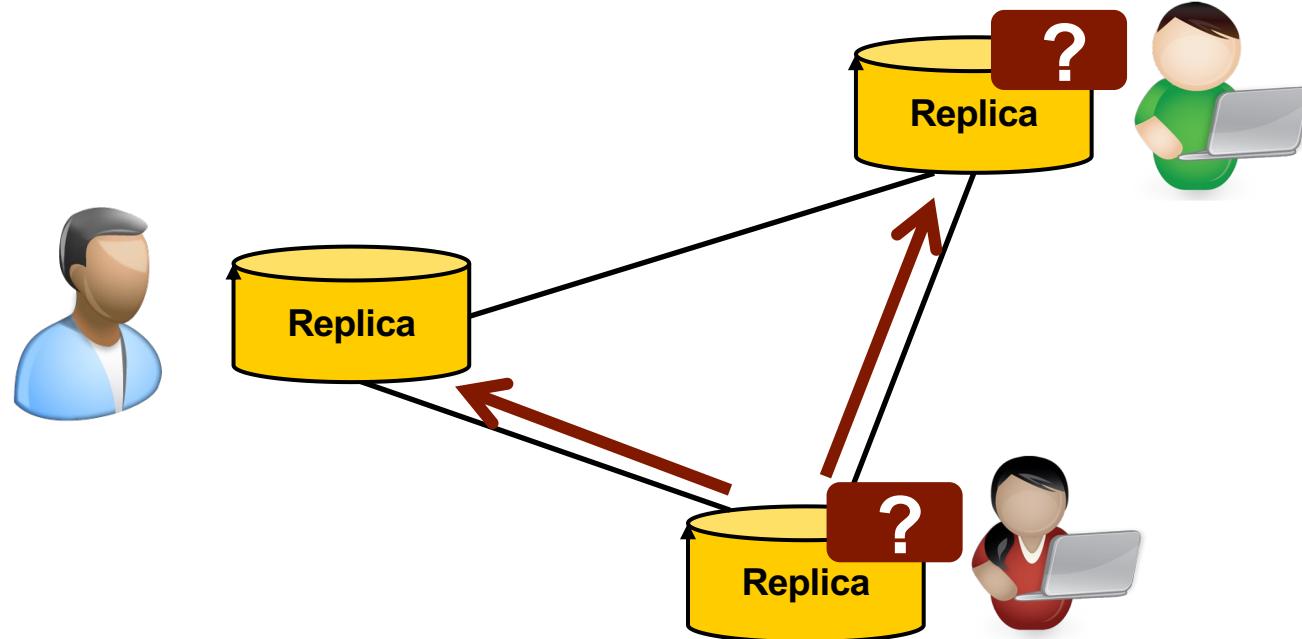
Example: Replicated Data



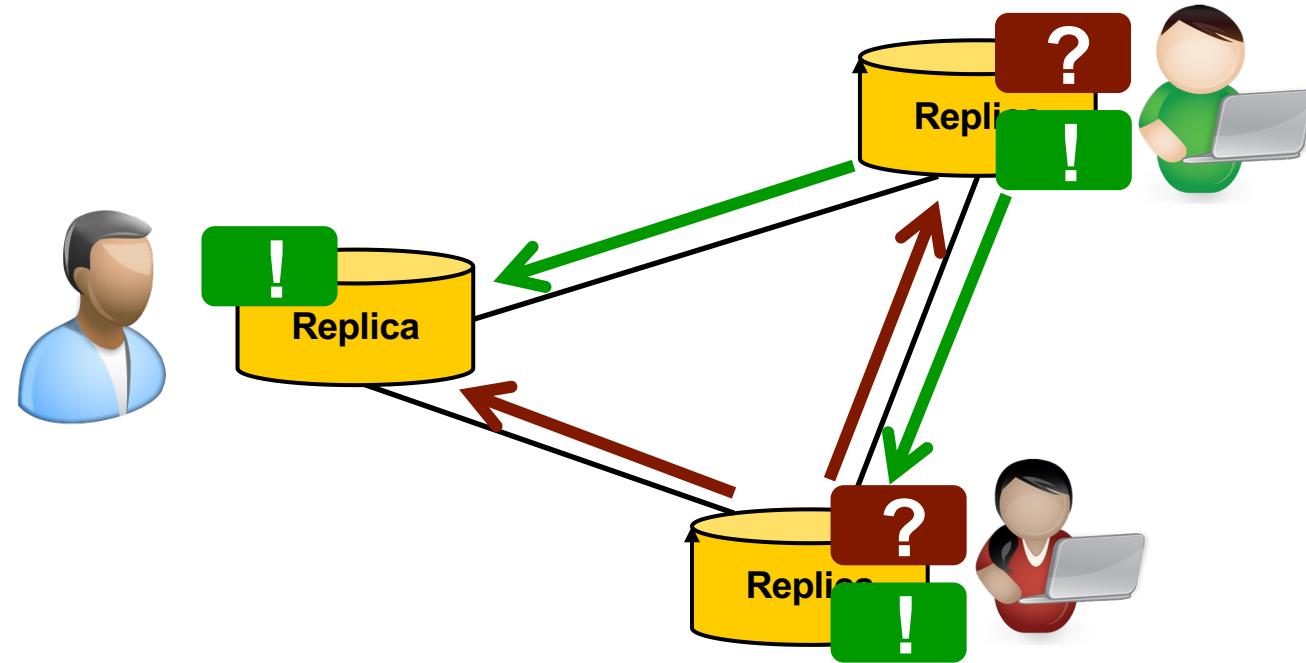
Example: Consistency



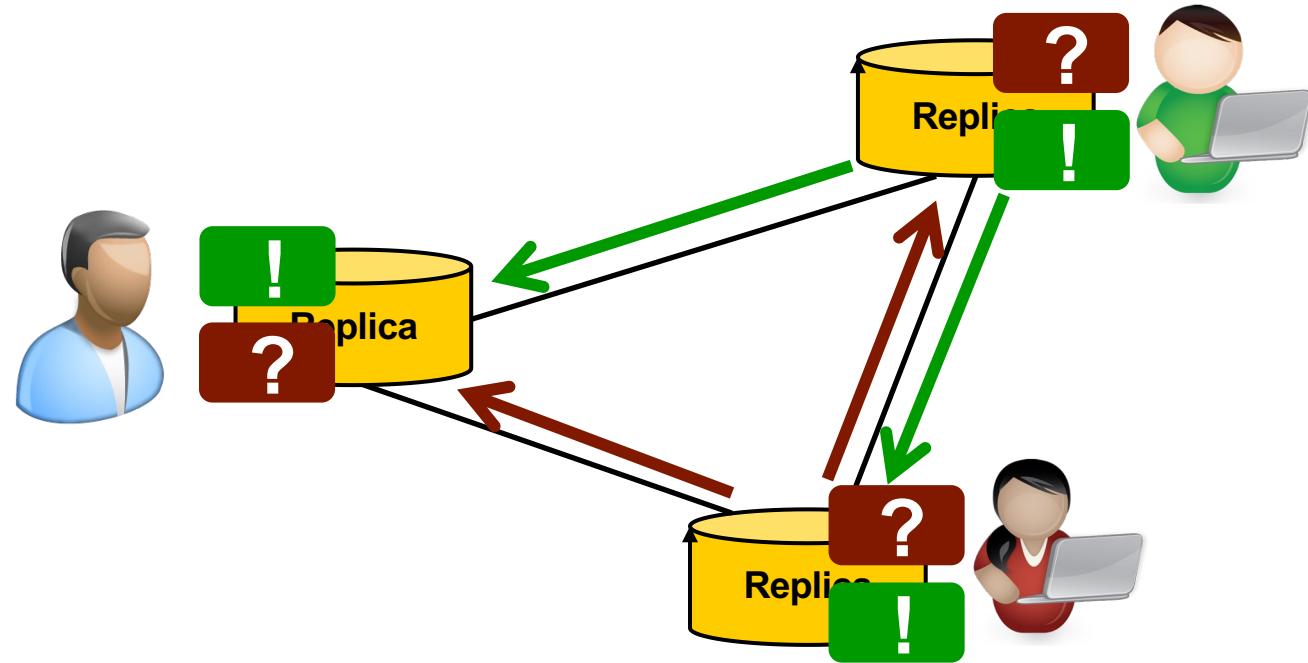
Example: Consistency



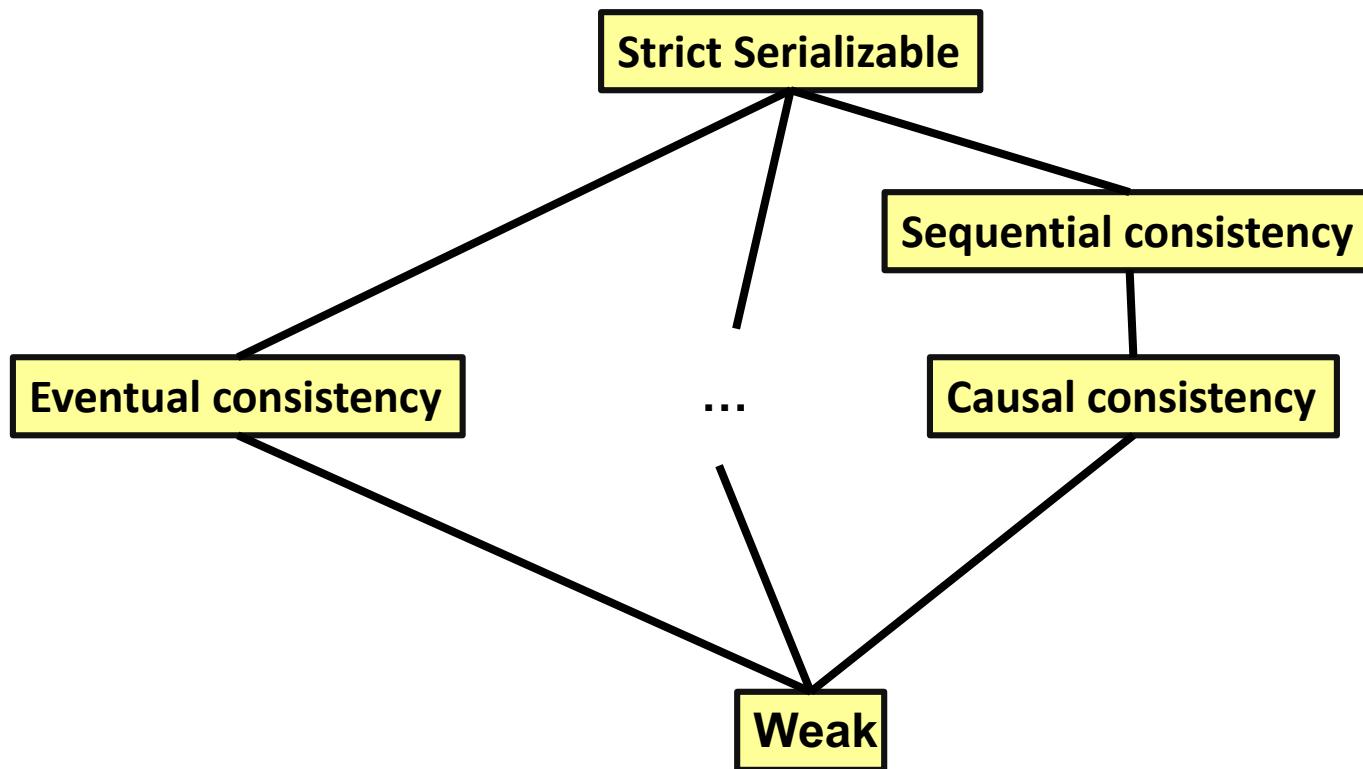
Example: Consistency



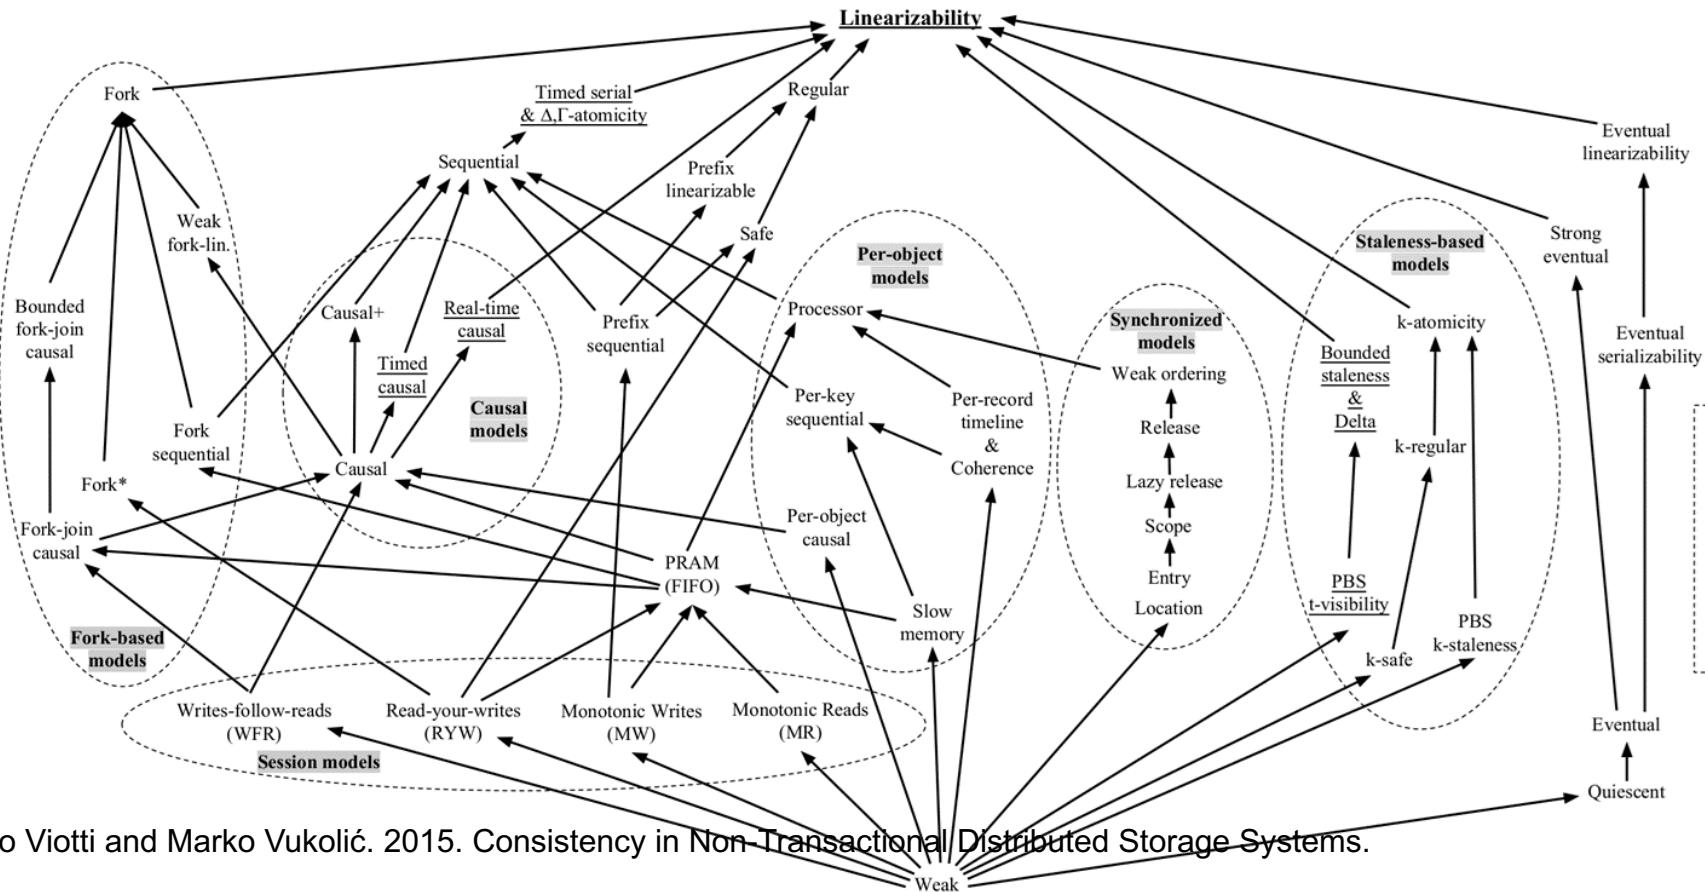
Example: Consistency



Consistency Levels



Consistency Levels



Paolo Viotti and Marko Vukolić. 2015. Consistency in Non-Transactional Distributed Storage Systems.

Eventual consistency

Causal consistency

Fifo consistency

...

DATASTAX

Apache Cassandra™ 2.1 (Supported)

Cassandra 2.1 (used by DSE 4.7, 4.8) ▾

About Cassandra What's new CQL Understanding the architecture Planning a deployment Installing
Initializing a cluster Security Database internals Configuration Operations Backing up and restoring data
Cassandra tools References Moving data to/from other databases Troubleshooting Release notes

Search Advanced search

Home / Database internals / Data consistency / Configuring data consistency

Configuring data consistency

Consistency levels in Cassandra can be configured to manage availability versus data accuracy. You can configure consistency on a cluster, datacenter, or individual I/O operation basis. Consistency among participating nodes can be set globally and also controlled on a per-operation basis (for example insert or update) using Cassandra's drivers and client libraries.

Write consistency levels

This table describes the write consistency levels in strongest-to-weakest order.

Write Consistency Levels

Level	Description	Usage
ALL	A write must be written to the commit log and memtable on all replica nodes in the cluster for that partition.	Provides the highest consistency and the lowest availability of any other level.
EACH_Q_UORUM	Strong consistency. A write must be written to the commit log and memtable on a quorum of replica nodes in <i>all</i> datacenter.	Used in multiple datacenter clusters to strictly maintain consistency at the same level in each datacenter. For example, choose this level if you want a read to fail when a datacenter is

Menu  English Sign In to the Console

Amazon DynamoDB Developer Guide (API Version 2012-08-10)

Documentation - This Guide Search

AWS Documentation » Amazon DynamoDB » Developer Guide » What Is Amazon DynamoDB? » Amazon DynamoDB: How It Works » Read Consistency

Read Consistency

Amazon DynamoDB is available in multiple AWS regions around the world. Each region is completely independent and isolated from other AWS regions. For example, if you have a table called `People` in the `us-east-1` region and another table named `People` in the `us-west-2` region, these are considered two entirely separate tables. For a list of all the AWS regions in which DynamoDB is available, see [AWS Regions and Endpoints](#) in the *Amazon Web Services General Reference*.

Every AWS region consists of multiple distinct locations called Availability Zones. Each Availability Zone is isolated from failures in other Availability Zones, and to provide inexpensive, low-latency network connectivity to other Availability Zones in the same region. This allows rapid replication of your data among multiple Availability Zones in a region.

When your application writes data to a DynamoDB table and receives an HTTP 200 response (OK), all copies of the data are updated. The data will eventually be consistent across all storage locations, usually within one second or less.

DynamoDB supports *eventually consistent* and *strongly consistent* reads.

Eventually Consistent Reads

When you read data from a DynamoDB table, the response might not reflect the results of a recently completed write operation. The response might include some stale data. If you repeat your read request after a short time, the response should return the latest data.

Strongly Consistent Reads

When you request a strongly consistent read, DynamoDB returns a response with the most up-to-date data, reflecting the updates from all prior write operations that were successful. A strongly consistent read might not be available in the case of a network delay or outage.

Note

DynamoDB uses eventually consistent reads, unless you specify otherwise. Read

rame: During the

mongoDB | FOR GIANT IDEAS

How does MongoDB ensure consistency?

Back to Table of Contents

MongoDB is consistent by default: reads and writes are issued to the primary member of a replica set. Applications can optionally read from secondary replicas, where data is eventually consistent by default. Reads from secondaries can be useful in scenarios where it is acceptable for data to be slightly out of date, such as some reporting applications. Applications can also read from the closest copy of the data (as measured by ping distance) when latency is more important than consistency.

Learn more in the [MongoDB Architecture Guide](#).

ConSysT Programming Framework

Combine multiple consistency levels in the same application

High consistency (slow)

Payment, Security, ...

Low consistency (fast)

Logging, messaging, ...



ConSysT Programming Framework

```
class Concert {  
    Date date;  
    Ref<@Weak ConcertHall> hall;  
    Ref<@Weak Band> band;  
    Ref<@Strong Counter> soldTickets;  
  
    ...  
}
```

Programming model supports explicit consistency levels

Information flow static analysis ensures safe mixing of consistency levels



Type System: Ensures Safety

```
Ref<@Weak Counter> weakCounter;  
Ref<@Strong Counter> strongCounter;
```

Allowed

```
weakCounter.ref.value = strongCounter.ref.value;
```

Strong to Weak

```
Ref<@Weak Counter> weakCounter;  
Ref<@Strong Counter> strongCounter;
```

Disallowed

```
strongCounter.ref.value = weakCounter.ref.value;
```

Weak to Strong

```
Ref<@Weak Counter> weakCounter;  
Ref<@Strong Counter> strongCounter;
```

Disallowed

```
if (weakCounter.ref.value == 0) {  
    strongCounter.ref.value = 0;  
}
```

Weak to Strong (implicit)

ConSysT

Tunable, safe consistency meets object-oriented programming.

<https://consyst-project.github.io/>

What is ConSysT?

ConSysT is a distributed object-oriented language. Objects can be replicated with different levels of **consistency**. The type system ensures that consistency levels are mixed safely.

Multiple consistency levels

Each replicated object comes with its own consistency level.

Safe mixing of consistencies

The static type system ensures correct mixing of consistency levels.

Object-oriented programming

Consistency is fully integrated with object-oriented abstraction.

Rethinking Safe Consistency in Distributed Object-Oriented Programming

MIRKO KÖHLER, Technische Universität Darmstadt, Germany

NAFISE ESKANDANI, Technische Universität Darmstadt, Germany

PASCAL WEISENBURGER, Technische Universität Darmstadt, Germany

ALESSANDRO MARGARA, Politecnico di Milano, Italy

GUIDO SALVANESCHI, Universität St. Gallen, Switzerland

Large scale distributed systems require to embrace the trade off between consistency and availability, accepting lower levels of consistency to guarantee higher availability. Existing programming languages are, however, agnostic to this compromise, resulting in consistency guarantees that are the same for the whole application and are implicitly adopted from the middleware or hardcoded in configuration files.

In this paper, we propose to integrate availability in the design of an object-oriented language, allowing developers to specify different consistency and isolation constraints in the same application at the granularity of single objects. We investigate how availability levels interact with object structure and define a type system that preserves correct program behavior. Our evaluation shows that our solution performs efficiently and improves the design of distributed applications.

CCS Concepts: • Software and its engineering → Distributed programming languages; • Information systems → Remote replication.

Additional Key Words and Phrases: replication, consistency, type systems, Java

ACM Reference Format:

Mirko Köhler, Nafise Eskandani, Pascal Weisenburger, Alessandro Margara, and Guido Salvaneschi. 2020. Rethinking Safe Consistency in Distributed Object-Oriented Programming. *Proc. ACM Program. Lang.* 4, OOPSLA, Article 188 (November 2020), 30 pages. <https://doi.org/10.1145/3428256>

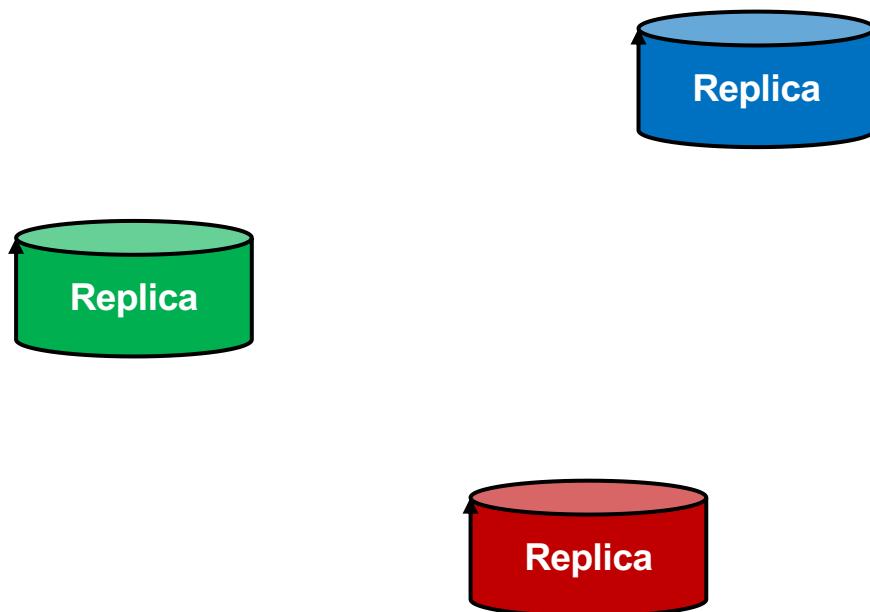
Overview

In **ConSysT**, the main abstraction are *replicated objects* that are fully integrated into an *object-oriented* language. Replicated objects have a *consistency level* specified by the developer.

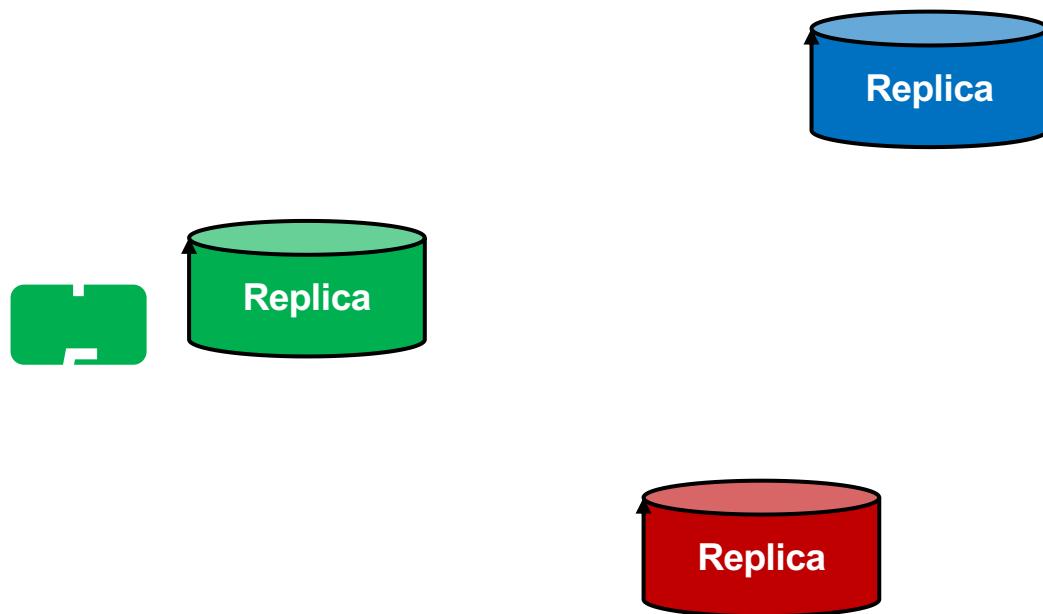
Distribution

Easily distribute your data across your local network, datacenters or geo-replicated devices. *Replicated objects* allow to distribute and perform operations on your data. ConSysT is implemented as a language extension to Java, you can create replicated

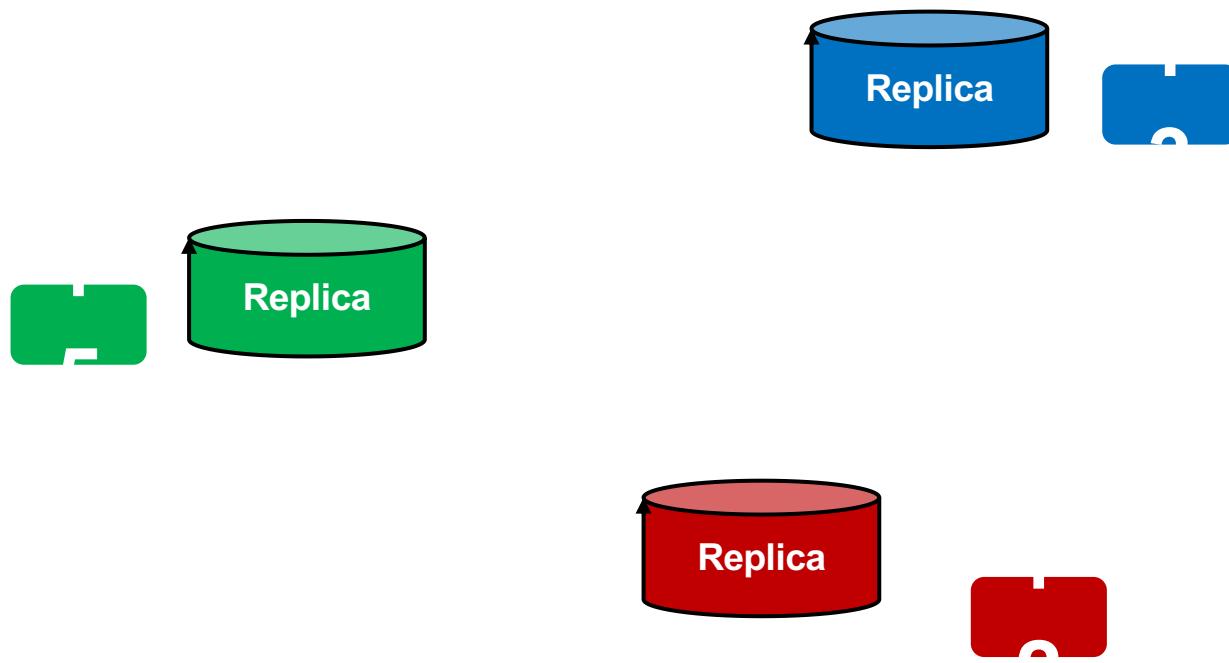
Eventual Consistency: CRDTs



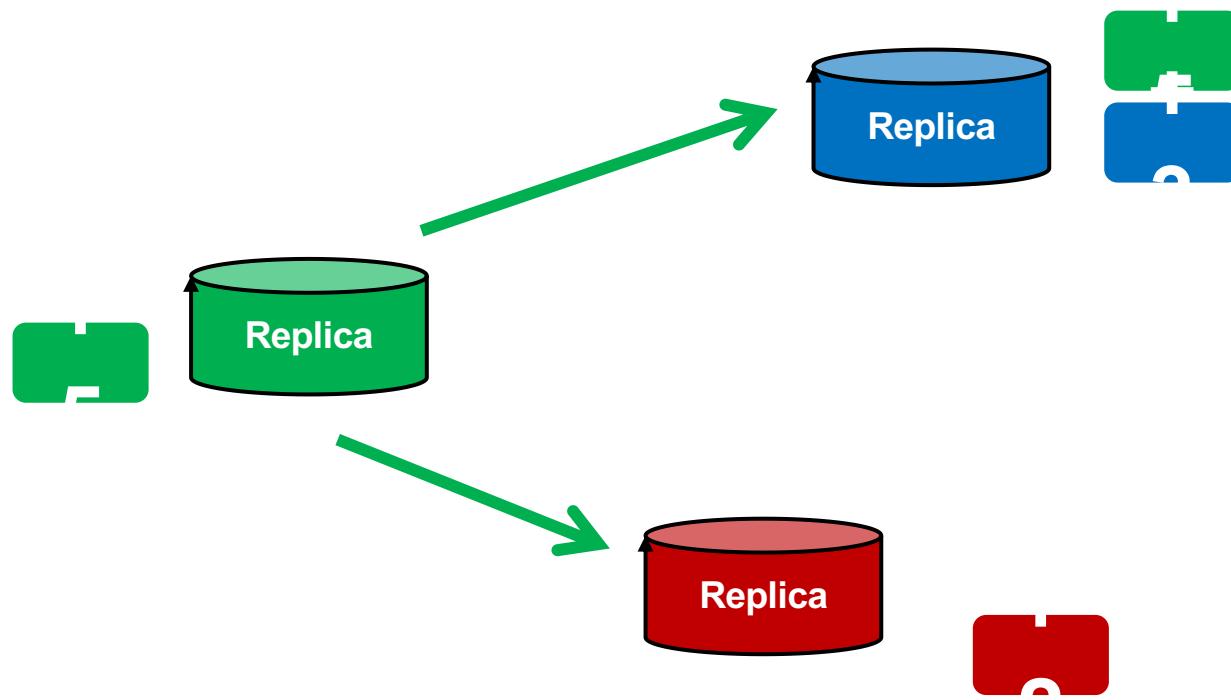
Eventual Consistency: CRDTs



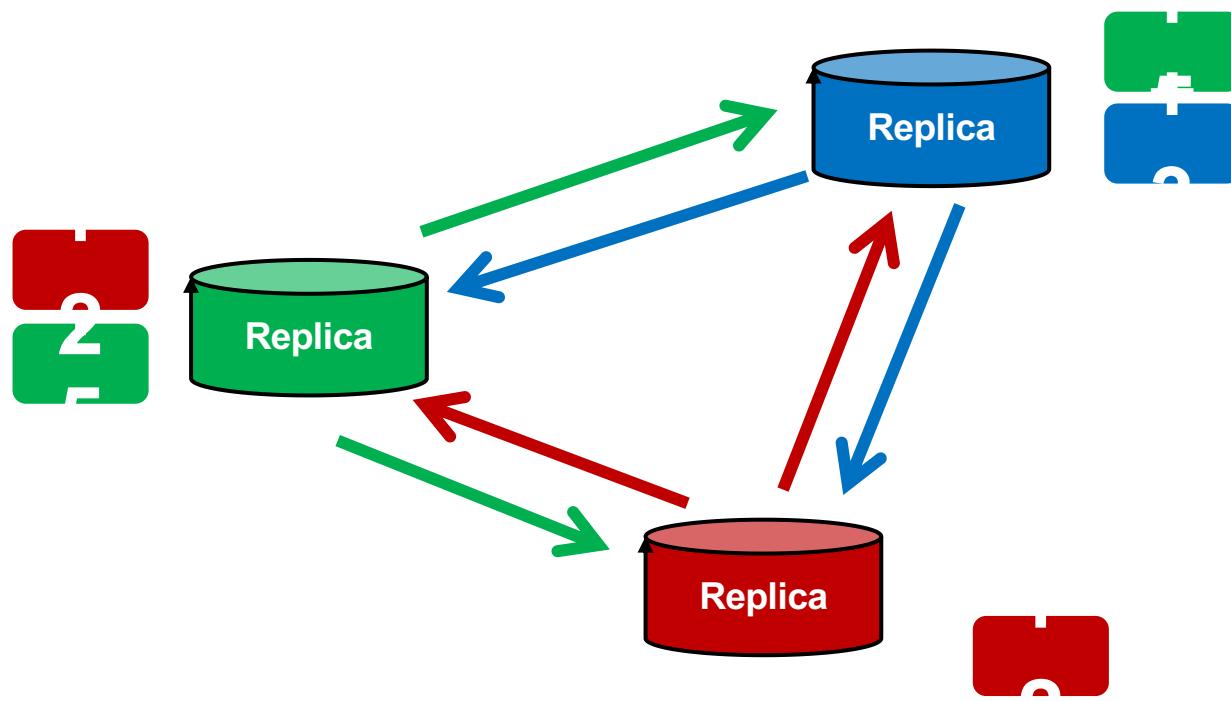
Eventual Consistency: CRDTs



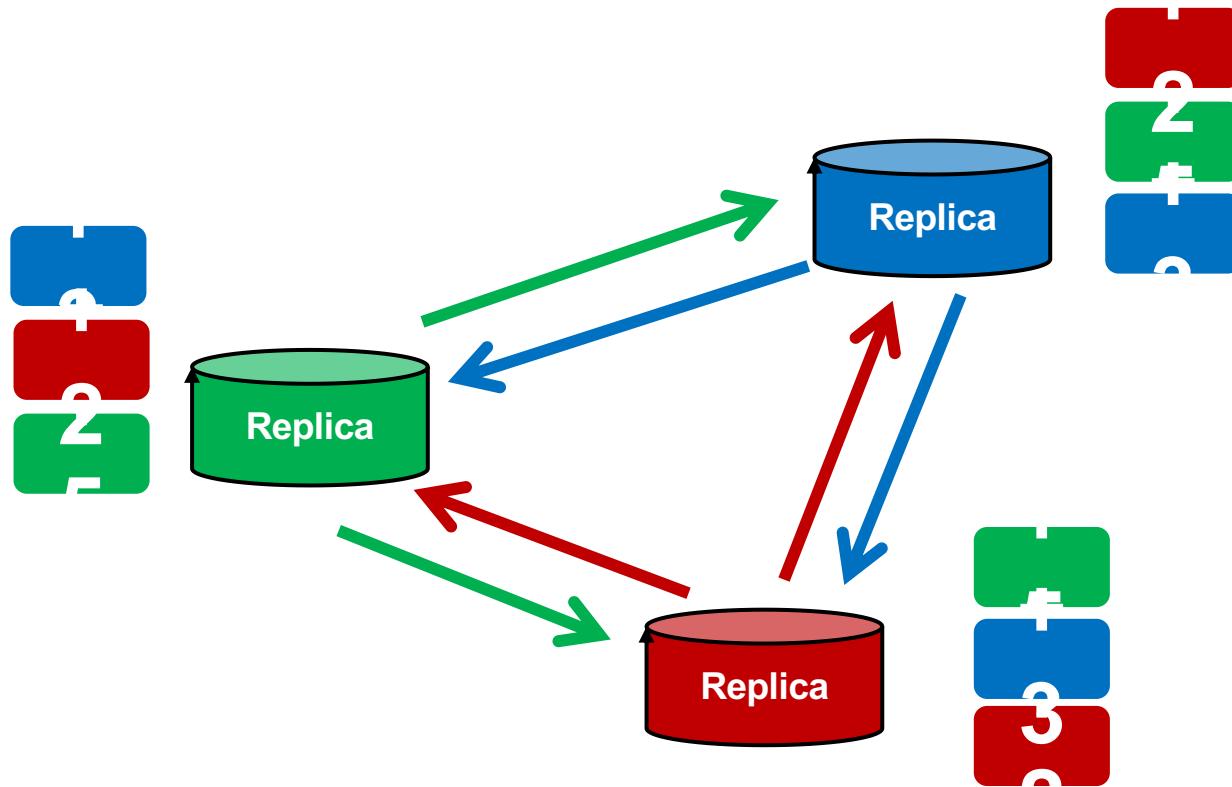
Eventual Consistency: CRDTs



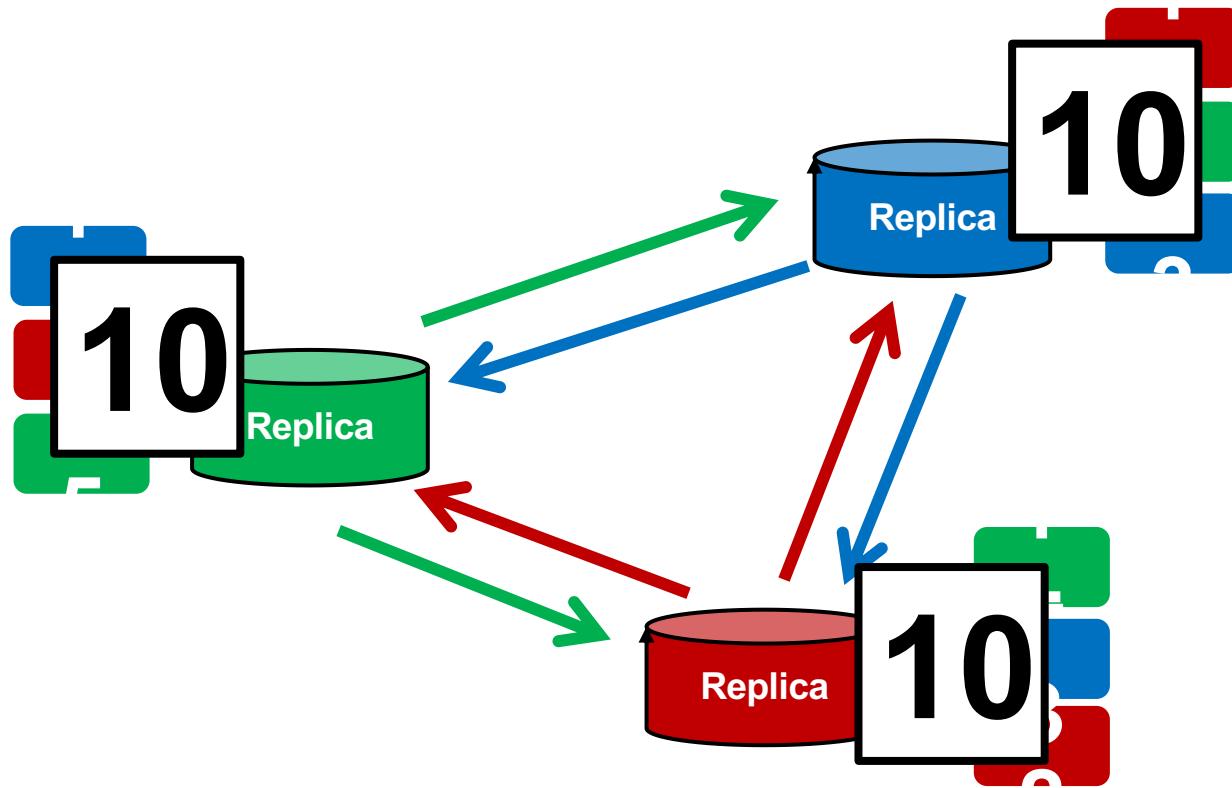
Eventual Consistency: CRDTs



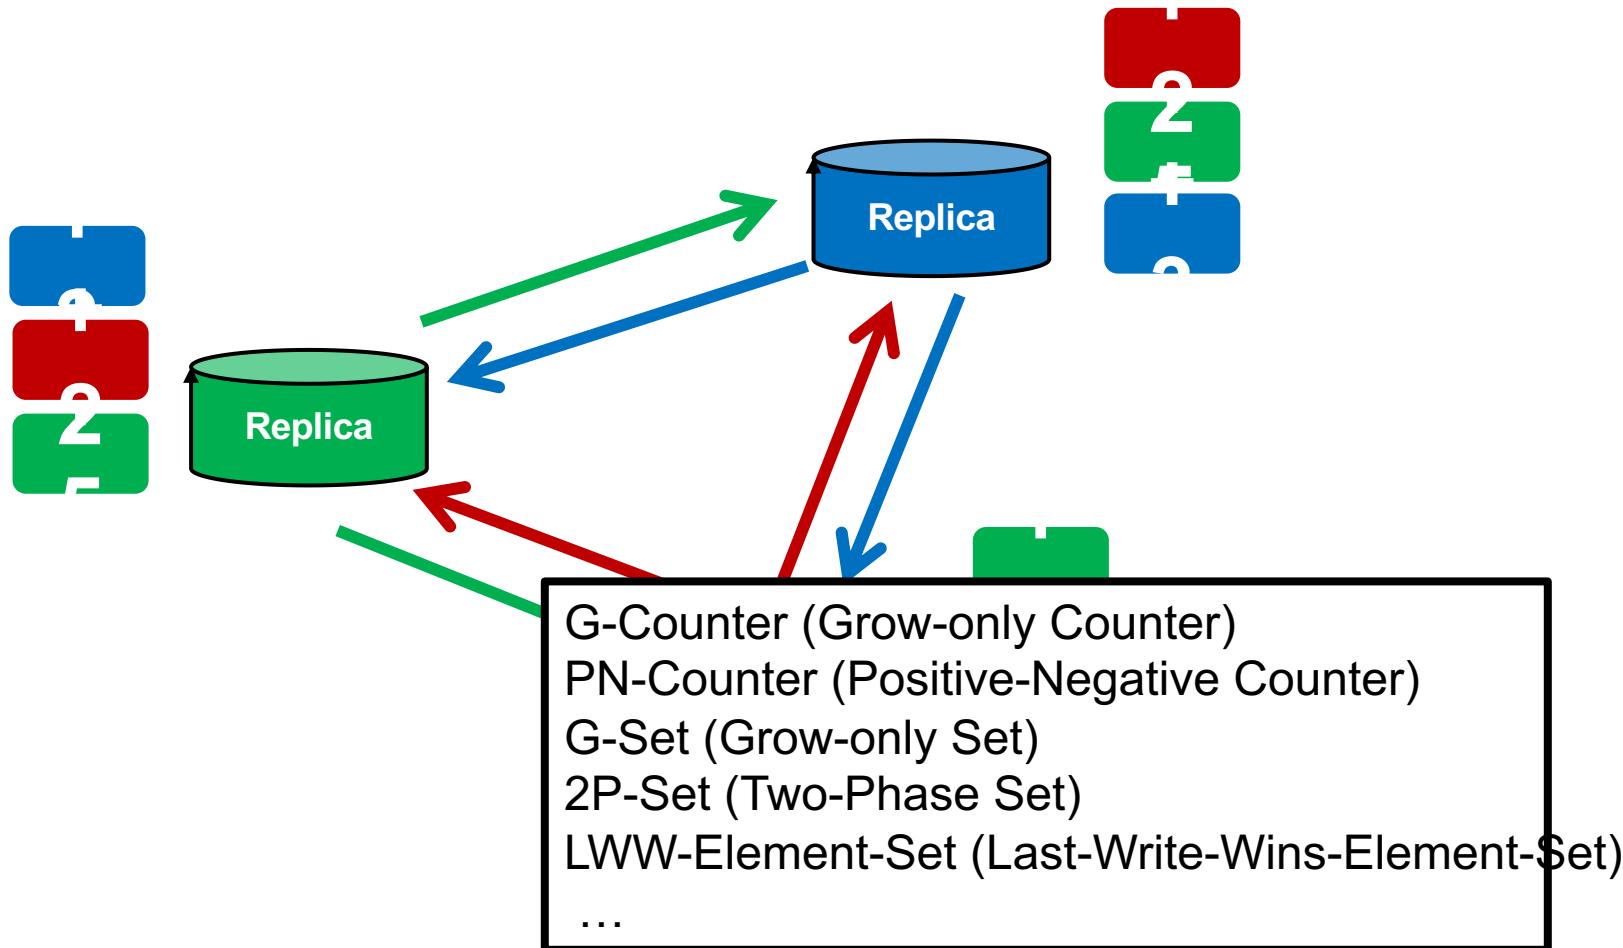
Eventual Consistency: CRDTs



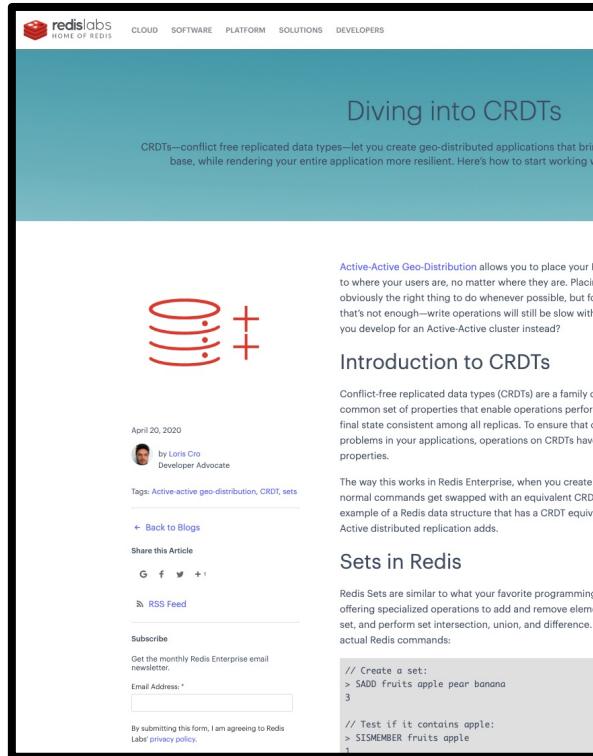
Eventual Consistency: CRDTs



Eventual Consistency: CRDTs



Eventual Consistency: CRDTs



The screenshot shows a blog post titled 'Diving into CRDTs' on the Redis Labs website. The post discusses conflict-free replicated data types (CRDTs) and their use in geo-distributed applications. It features a red icon of a stack of three circles with a plus sign. The author is Lora Cro, a Developer Advocate. The post is dated April 20, 2020. It includes sections on 'Introduction to CRDTs' and 'Sets in Redis', with Redis commands for creating sets and testing membership.

```
// Create a set:  
> SADD fruits apple pear banana  
3  
  
// Test if it contains apple:  
> SISMEMBER fruits apple  
1
```



The screenshot shows the AntidoteDB landing page. The main headline is 'A cure for consistency under partitioning'. It describes AntidoteDB as a synchronization-free execution system that provides fault-tolerance and high availability. A 'GET STARTED!' button is present. The page is divided into sections: 'Features' (with icons for CRDTs, Highly Available Transactions, and Geo-replication), and a 'LWW-Element-Set (Last-Write-Wins-Element-Set)' section at the bottom.

A cure for consistency under partitioning

AntidoteDB operations is based on the principle of synchronization-free execution. Similarly to NOSQL, it provides fault-tolerance and high availability, but unlike NOSQL, AntidoteDB provides highly available SQL transactions that makes development easy.

GET STARTED!

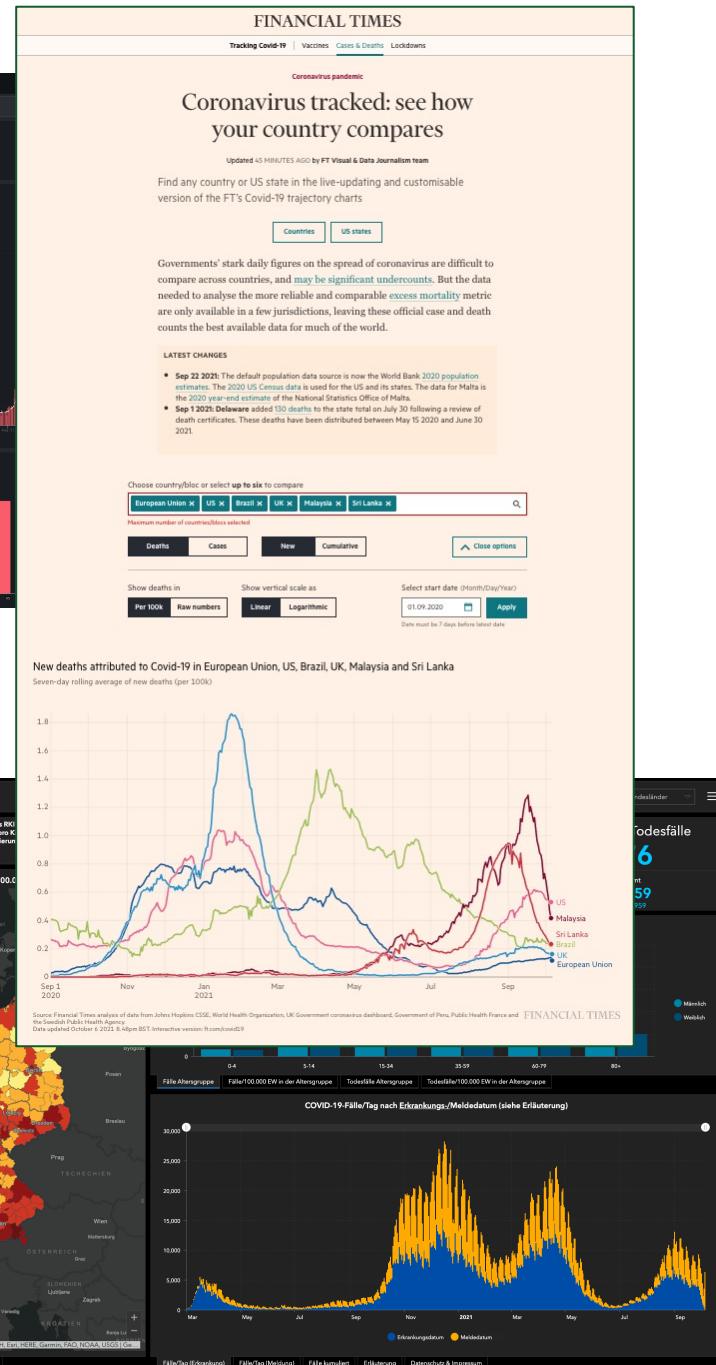
Features

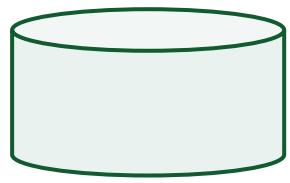
- CRDTs**
High-level replicated data types that are designed to work correctly in the presence of concurrent updates and partial failures.
- Highly Available Transactions**
Traditional ACID transactions were built for single-machine deployments. It is expensive to implement ACID transactions in distributed deployments. On the other hand, highly-available transactions (HAT) provide strong consistency within a data center, but still perform well in geo-replicated deployments.
- Geo-replication**
Designed to run on multiple servers in locations distributed world-wide. It provides continuous functioning even when there are failures or network partition.

LWW-Element-Set (Last-Write-Wins-Element-Set)

...

WHAT'S NEXT?





DATA



AGGREGATION
FILTERING, ...



GRAPHICS



GLOBAL PROGRAMMING



DATA



AGGREGATION
FILTERING, ...



GRAPHICS

Secure Enclave Programming

Enclaves: hardware protection for data and code from system software

- ~250MB memory
- Native speed

Here to stay

- Google Cloud announces to offer enclaves to customers (Summer '20)
- Intel announces 3rd Gen Xeon Processors: enclaves up to 1TB (Oct '20)



```

1 enclave {
2     trusted {
3         public void checkPassword([in, size=len] char* guess,
4             [out] int* result, size_t len);
5 }

```

enclave.edl

```

1 const char* password = "secret";
2 void checkPassword(char* guess, int* result, size_t len) {
3     strcmp(guess, password) == 0) ?
4         *result = 1 : *result = 0;
5 }

```

enclave.cpp

```

1 #include "sgx_urts.h"
2 #include "enclave_u.h"
3 #define BUF_LEN 100
4 int main() {
5     sgx_enclave_id_t eid;
6     sgx_status_t ret = SGX_SUCCESS;
7     sgx_launch_token_t token = {0};
8     int updated = 0;
9     ret = sgx_create_enclave(ENCLAVE_FILE, SGX_DEBUG_FLAG,
10         &token, &updated, &eid, NULL);
11     if (ret != SGX_SUCCESS) { ... /* exception */ }
12     char* guess = ... // read guess from stdin
13     int result = 0;
14     checkPassword(eid, guess, &result, BUF_LEN);
15     if (SGX_SUCCESS != sgx_destroy_enclave(eid)) {...}
16     return 0;
}

```

main.cpp

```

1 @Enclave
2 class PasswordChecker {
3     @Secret static String password = "secret";
4
5     @Gateway
6     public static boolean checkPassword(String guess) {
7         String guessE = endorse(guess);
8         return declassify(guessE.equals(password));
9     }
}

```

PasswordChecker.java

```

1 class Main {
2     public static void main(String[] args) {
3         String guess = ... // read guess from stdin
4         PasswordChecker.checkPassword(guess);
5     }
}

```

Main.java

JE*

C++ - Intel Visual Studio

[A. Oak, A. M. Ahmadian, M. Balliu, G. Salvaneschi, **Language Support for Secure Software Development with Enclaves, CSF'21**]

(Modeling) Infrastructure as Code

Example:

- Create a Web server in Google with TypeScript on Node.js.
- Create the instance, export the IP and Hostname.

Reactivity to external events

Deployment safety?

Deployment patterns?

[Automating Serverless Deployments for DevOps Organizations, D. Sokolowski, P. Weisenburger, G. Salvaneschi, FSE'21]

```
const gcp = require("@pulumi/gcp");

const computeNetwork = new gcp.compute.Network("network", {
  autoCreateSubnetworks: true,
});

const computeFirewall = new gcp.compute.Firewall("firewall", {
  network: computeNetwork.selfLink,
  allows: [
    {
      protocol: "tcp",
      ports: [ "22", "80" ],
    },
  ],
});

// Create a simple web server.
const startupScript =
  "#!/bin/bash \n" +
  "echo 'Hello, World!' > index.html \n" +
  "nohup python -m SimpleHTTPServer 80 &";

const computeInstance = new gcp.compute.Instance("instance", {
  machineType: "f1-micro",
  metadataStartupScript: startupScript,
  bootDisk: {
    initializeParams: {
      image: "debian-cloud/debian-8",
    },
  },
  networkInterfaces: [
    {
      network: computeNetwork.id,
      accessConfigs: [{}], // must be empty
    },
  ],
  serviceAccount: {
    scopes: ["https://www.googleapis.com/auth/cloud-platform"],
  },
}, { dependsOn: [computeFirewall] });

exports.instanceName = computeInstance.name;
exports.instanceIP = computeInstance.networkInterfaces.apply(ni => ni[0].ipAddress);
```

Smart Contracts

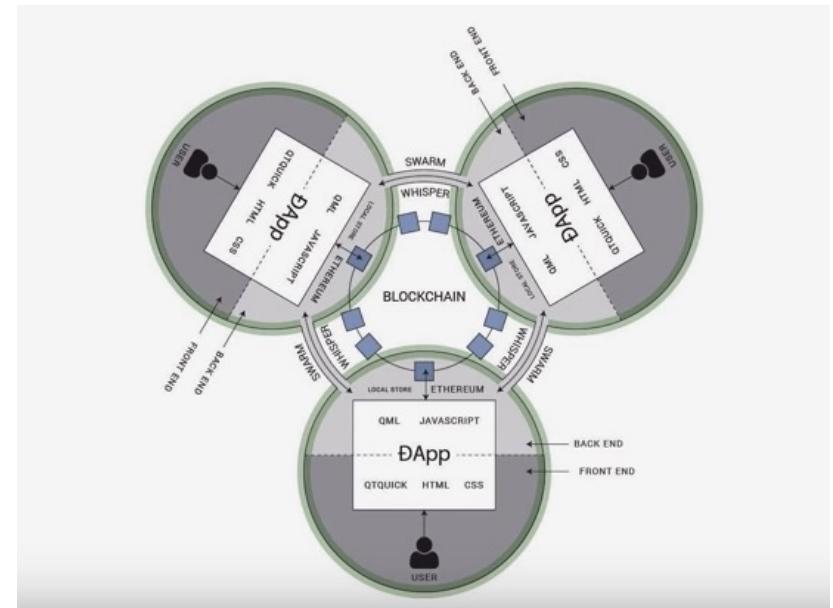
Single language for clients and contract

Client compiles to JavaScript

Contract compiles to Solidity

Multitier
programming?

val gui on Client
val score on Contract





Thank you

THANKS!



@guidosalva

www.guidosalvaneschi.com

Barbara Liskov

Carlo Ghezzi

Mira Mezini

Andrew Myers

Patrick Eugster

Alessandro Margara

Nada Amin

Philipp Haller

Ragnar Mogk

Sebastian Erdweg

Joscha Drechsler

Sven Amann

Ralf Mitschke

Sebastian Proksch

Sarah Nadi